



CAPSTONE PROJECT

PREDICTING CUSTOMER EVENT RESPONSE USING
CLASSIFICATION MODELS FOR A
CUSTOMER-TRANSACTIONS PROJECT



Contents

1. Introduction

2. Dataset

3. Insights

4. Data
Pre-processing

5. Exploratory
Data Analysis

6. Feature
Selection

7. Modelling

8. Visualization

9. Conclusion

Introduction

The Customer-Transactions project is a data visualization and analytics initiative that utilizes a sample dataset that simulates customer transaction data for a fictional office supplies store, covering sales, profit, orders, and customers.



Dataset

The given dataset consists of two Excel files named 'Customer by Store' and 'Transactions'

2851 Rows in Customer by Store
9426 Rows in Transactions

Customer by Store

Customer ID
Store Number
Customer Segment
Responder
Name
Address
Postcode
CustomerType
AverageTransaction
LastTransactionAmount
TransactionsPerYear
Income
HomeValue
Lat
Lon

Transactions

Customer ID
Order_ID
Order_Priority
Discount
Unit_Price
Quantity_ordered_new
Order_Date
Ship_Date
Shipping_Cost
Ship_Mode
Product_Category
Product_Sub-Category
Product_Container
Product_Base_Margin
Profit

Insights

The manager from the sales department would like to gain insights from the data. Some of the questions are as follows:

- What is the profit by customer segment and store no.?
- Which store has the highest customer?
- How many transactions do not match with customer info?
- Estimate (predict) the responder (Yes/No) based on the customer segment and another appropriate input
- What is the summarised value (Total) for Average Transaction, Last Transaction and Transaction Per Year by customer type?



Data Preprocessing



Load Data

```
1 # read the excel file and put it into CustomerStore dataframe
2 CustomerStore = pd.ExcelFile("Customers by Store.xlsx")
3
4 # List all sheet names
5 sheet_names = CustomerStore.sheet_names
6
7 # Read all sheets into a list of DataFrames
8 CustomerByStore = [pd.read_excel(CustomerStore, sheet_name=sheet) for sheet in
9 sheet_names]
10
11 # Concatenate all DataFrames into one
12 combined_df = pd.concat(CustomerByStore, ignore_index=True)
13
14 # Save the combined DataFrame to a new Excel file (in Google Colab environment)
15 combined_df.to_excel('combined_output.xlsx', index=False)
```

```
1 Transactions = pd.read_excel("transactions.xlsx")
2 Transactions.head()
```



Merge Data

```
1 # Load the two Excel files
2 file1 = 'transactions.xlsx'
3 file2 = 'combined_output.xlsx'
4
5 # Read the specific sheets or the first sheet of each file
6 Transactions = pd.read_excel(file1, sheet_name='sheet1')
7 combined_df = pd.read_excel(file2, sheet_name='Sheet1')
8
9 # Merge the two DataFrames based on 'Customer ID' using inner join. This is to
10 # make sure that only the data with a valid Customer ID will be merged.
11 mergedfile_df = pd.merge(Transactions, combined_df, left_on='Customer_ID',
12 right_on='Customer ID', how='inner')
```


Data Preprocessing



Clean Data

```
1 # Check missing values, are all values 0?
2 merged_df.isna().sum()
```

```
1 # Remove rows with missing values.
2 merged_df.dropna(inplace=True)
```

```
1 CustTransaction.info()
```

```
1 # Convert Order_Date to datetime (assuming the format is mmddyyyy)
2 CustTransaction['Order_Date'] = pd.to_datetime(CustTransaction['Order_Date'])
3 CustTransaction['Ship_Date'] = pd.to_datetime(CustTransaction['Ship_Date'])
```

```
1 from sklearn.preprocessing import LabelEncoder
2 encoder = LabelEncoder()
3
4 # Select all the columns with object datatype that need to be encoded
5 col = ['Order_Priority', 'Ship_Mode', 'Product_Category',
        'Product_Sub-Category', 'Product_Container', 'Customer Segment', 'Responder',
        'Postcode', 'CustomerType']
6
7 # Transform the data
8 for i in col:
9     CustTransactionClassification[i] = encoder.fit_transform
        (CustTransactionClassification[i])
```



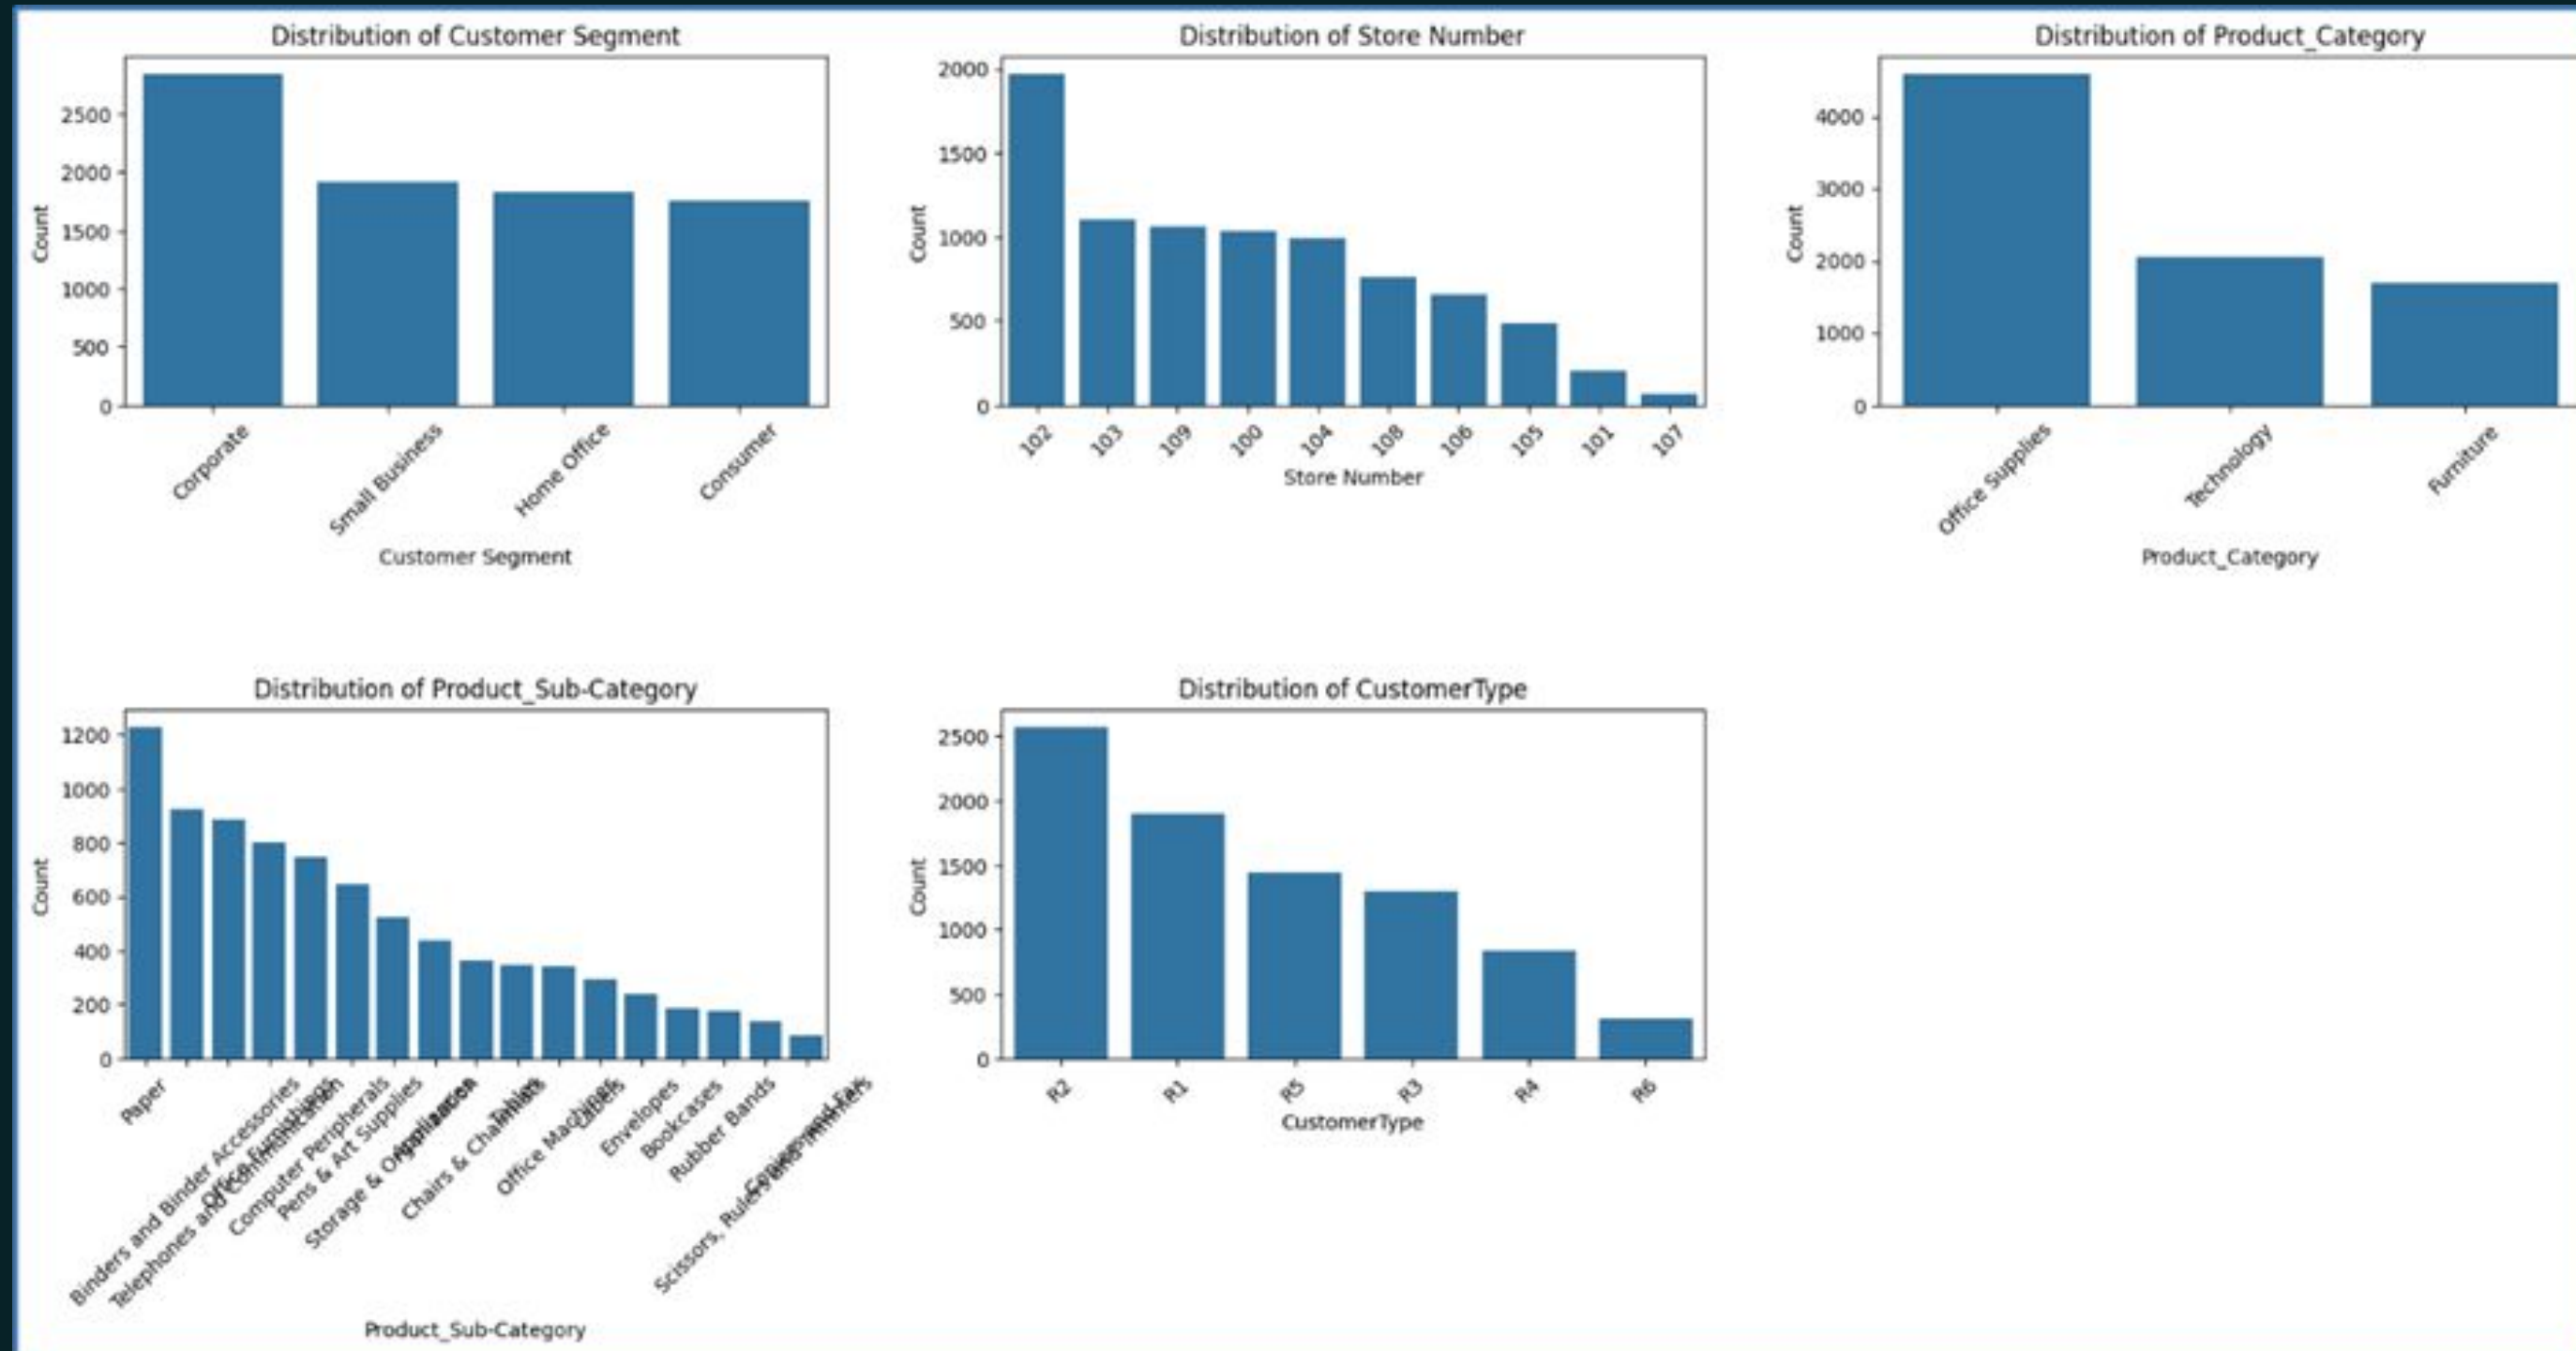
Export Data

```
1 CustTransaction.to_excel('CustTransaction.xlsx', index=False)
```

```
1 from google.colab import files
2 files.download('CustTransaction.xlsx')
```

Exploratory Data Analysis

Data Distribution



Exploratory Data Analysis

Data Correlation - Spearman rank's correlation coefficient

	Discount	Unit_Price	Quantity_ordered_new	Shipping_Cost	Product_Base_Margin	Profit	AverageTransaction	LastTransactionAmount	TransactionsPerYear	Income	HomeValue
Discount	1	-0.006273707	-0.011591588	-0.00340045	-0.000576748	-0.064887861	0.005281703	0.005851087	0.004238613	-0.004911294	-0.006424278
Unit_Price	-0.006273707	1	-0.039479181	0.649984896	0.399350072	0.208505328	-0.009248684	-0.010134597	-0.015947754	-0.015425719	-0.005722735
Quantity_ordered_new	-0.011591588	-0.039479181	1	-0.03043601	-0.003447643	0.207969625	-0.002555293	-0.001511665	0.041255511	-0.020158532	-0.042787842
Shipping_Cost	-0.00340045	0.649984896	-0.03043601	1	0.29663404	-0.124863538	-0.006314271	-0.005851825	-0.006638293	-0.014284819	-0.009893534
Product_Base_Margin	-0.000576748	0.399350072	-0.003447643	0.29663404	1	-0.131848372	-0.000822717	-0.000370578	0.000445982	0.003993226	0.005169546
Profit	-0.064887861	0.208505328	0.207969625	-0.124863538	-0.131848372	1	-0.011219557	-0.010446186	-0.012317973	-0.007648621	-0.018788029
AverageTransaction	0.005281703	-0.009248684	-0.002555293	-0.006314271	-0.000822717	-0.011219557	1	0.925767913	0.153715444	0.04521585	0.171308516
LastTransactionAmount	0.005851087	-0.010134597	-0.001511665	-0.005851825	-0.000370578	-0.010446186	0.925767913	1	0.137948673	0.03808749	0.160062008
TransactionsPerYear	0.004238613	-0.015947754	0.041255511	-0.006638293	0.000445982	-0.012317973	0.153715444	0.137948673	1	-0.014380358	-0.001589061
Income	-0.004911294	-0.015425719	-0.020158532	-0.014284819	0.003993226	-0.007648621	0.04521585	0.03808749	-0.014380358	1	0.57905103
HomeValue	-0.006424278	-0.005722735	-0.042787842	-0.009893534	0.005169546	-0.018788029	0.171308516	0.160062008	-0.001589061	0.57905103	1

SIGNIFICANT CORRELATION: Shipping_Cost and Unit_Price
AverageTransaction and LastTransactionAmount
Income and HomeValue

Feature Selection

The variables that are more likely not suitable for ML are dropped out.

```
1 CustTransactionClassification = pd.DataFrame(CustTransaction.drop(columns=[  
    'Customer_ID', 'Order_ID', 'Order_Date', 'Ship_Date', 'Name', 'Address'])))
```

Another proposed feature selection is to remove Lat and Lon. Both of these feature sets will be tested through ML modelling to see which feature set will give better result in terms of model accuracy.

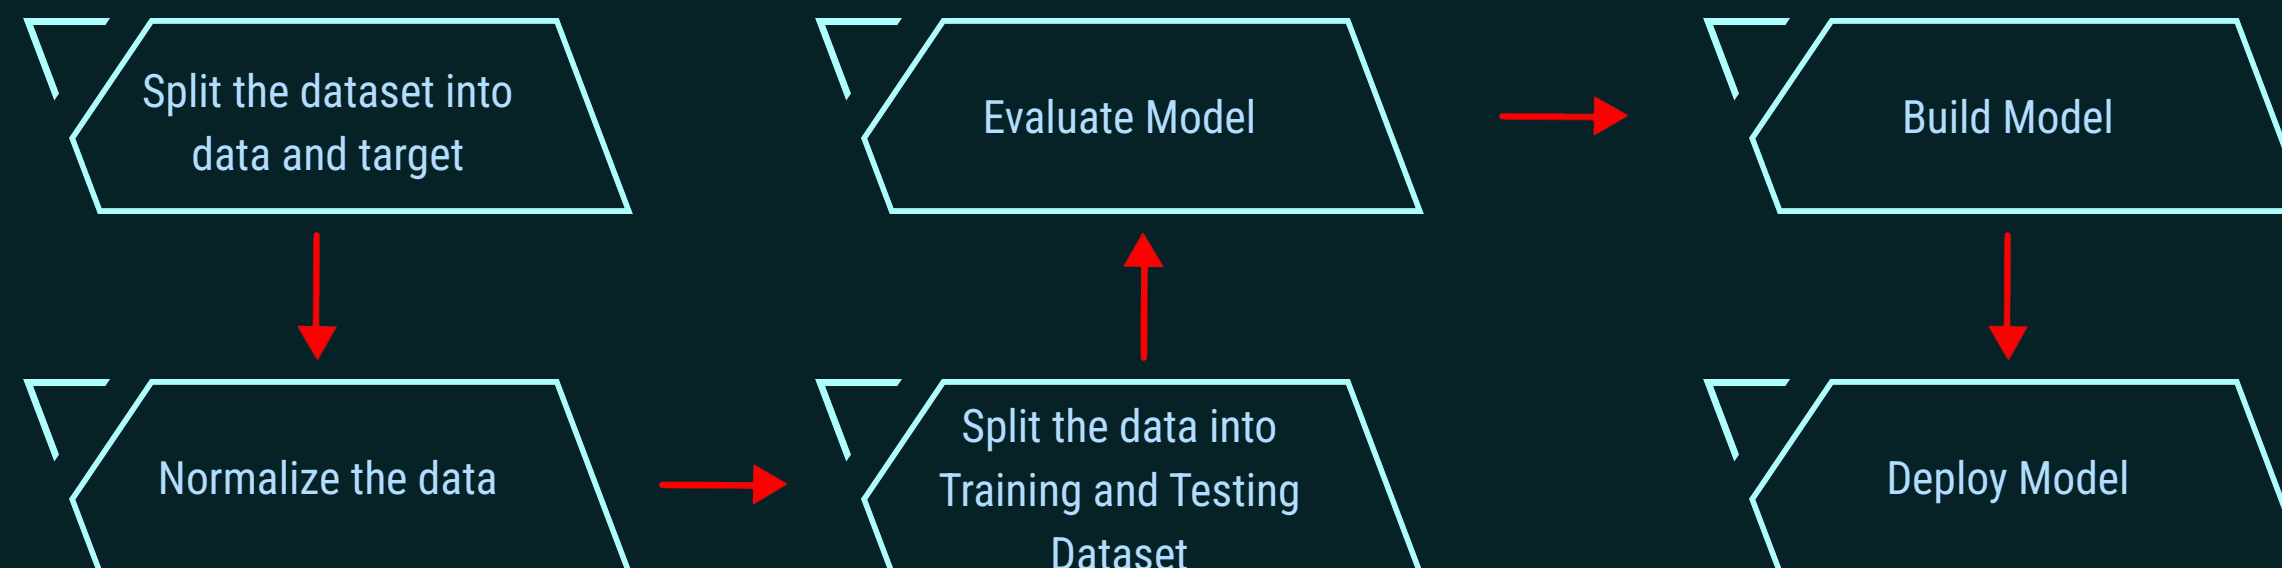
Selected Features

#	Column	Non-Null Count		Dtype
0	Order_Priority	8352	non-null	object
1	Discount	8352	non-null	float64
2	Unit_Price	8352	non-null	float64
3	Quantity_ordered_new	8352	non-null	int64
4	Shipping_Cost	8352	non-null	float64
5	Ship_Mode	8352	non-null	object
6	Product_Category	8352	non-null	object
7	Product_Sub-Category	8352	non-null	object
8	Product_Container	8352	non-null	object
9	Product_Base_Margin	8352	non-null	float64
10	Profit	8352	non-null	float64
11	Store Number	8352	non-null	int64
12	Customer Segment	8352	non-null	object
13	Responder	8352	non-null	object
14	Postcode	8352	non-null	object
15	CustomerType	8352	non-null	object
16	AverageTransaction	8352	non-null	float64
17	LastTransactionAmount	8352	non-null	float64
18	TransactionsPerYear	8352	non-null	float64
19	Income	8352	non-null	int64
20	HomeValue	8352	non-null	int64
21	Lat	8352	non-null	float64
22	Lon	8352	non-null	float64

Modelling

- The ML model focused in this capstone project is **CLASSIFICATION**, a type of prediction modeling used to categorize data into predefined classes or labels.
- This is to cater one of the insights to estimate (predict) the responder (Yes/No) based on the customer segment and another appropriate input
- To achieve accurate classification, three models were selected: K-Nearest Neighbors (KNN), Decision Tree, and Random Forest.

Modelling Steps



Modelling

1. Split the dataset into data and target

```
1 X = CustTransactionClassification.drop('Responder', axis=1)
2 Y = CustTransactionClassification['Responder']
```

2. Normalize the data

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4
5 # Normalize all the data in dataframe X using scaler and put it in dataframe
  Xscaled
6 Xscaled = scaler.fit_transform(X)
```

```
1 Xscaled
```


Modelling

3. Split into Training and Testing dataset

```
1 from sklearn.model_selection import train_test_split
2
3 xtrain, xtest, ytrain, ytest = train_test_split(Xscaled, Y, test_size=0.3,
    random_state=123)
```

K-Nearest Neighbours

```
1 # K-NEAREST NEIGHBOUR MODEL
2
3 from sklearn.neighbors import KNeighborsClassifier
4 knn_model = KNeighborsClassifier(n_neighbors=2)
5 knn_model.fit(xtrain, ytrain)
6 knn_prediction = knn_model.predict(xtest)
```

```
1 #EVALUATE K-NEAREST NEIGHBOUR MODEL
2
3 from sklearn import metrics
4 knn_CM = metrics.confusion_matrix(ytest, knn_prediction)
5 print('Confusion Matrix: \n', knn_CM)
6 print('\n')
7 knn_Acc = metrics.accuracy_score(ytest, knn_prediction)
8 print('Model accuracy is: ', knn_Acc)
```

Confusion Matrix:

```
[[1814   84]
 [ 431  177]]
```

Model accuracy is: 0.7944932162809257

Decision Tree

```
1  # DECISION TREE MODEL
2
3  from sklearn.tree import DecisionTreeClassifier
4  dt_model = DecisionTreeClassifier(criterion='entropy', random_state=123)
5  dt_model.fit(xtrain, ytrain)
6  dt_prediction = dt_model.predict(xtest)
```

```
1  #EVALUATE DECISION TREE MODEL
2
3  from sklearn import metrics
4  dt_CM = metrics.confusion_matrix(ytest, dt_prediction)
5  print('Confusion Matrix: \n', dt_CM)
6  print('\n')
7  dt_Acc = metrics.accuracy_score(ytest, dt_prediction)
8  print('Model accuracy is: ', dt_Acc)
```

Confusion Matrix:

```
[[1880   18]
 [  12 596]]
```

Model accuracy is: 0.9880287310454908

Random Forest

```
1  # RANDOM FOREST MODEL
2
3  from sklearn.ensemble import RandomForestClassifier
4  rf_model = RandomForestClassifier(n_estimators=300)
5  rf_model.fit(xtrain, ytrain) #TRAIN MODEL
6  rf_prediction = rf_model.predict(xtest) #TEST MODEL
```

```
1  #EVALUATE RANDOM FOREST MODEL
2
3  from sklearn import metrics
4  rf_CM = metrics.confusion_matrix(ytest, rf_prediction)
5  print('Confusion Matrix: \n', rf_CM)
6  print('\n')
7  rf_Acc = metrics.accuracy_score(ytest, rf_prediction)
8  print('Model accuracy is: ', rf_Acc)
```

Confusion Matrix:

```
[[1893   5]
 [  19 589]]
```

Model accuracy is: 0.9904229848363927

Modelling

Feature Set 1

(All the suitable features)

Classification Model	Accuracy
K-Nearest Neighbors	0.79449
Decision Tree	0.98803
Random Forest	0.99042
SVM (linear kernel)	0.78093
SVM (RBF kernel)	0.79888
SVM (sigmoid kernel)	0.70072
SVM (polynomial kernel)	0.79848
Neural Network Tensorflow	0.79409

Best Model: Random Forest (99%)

Feature Set 2

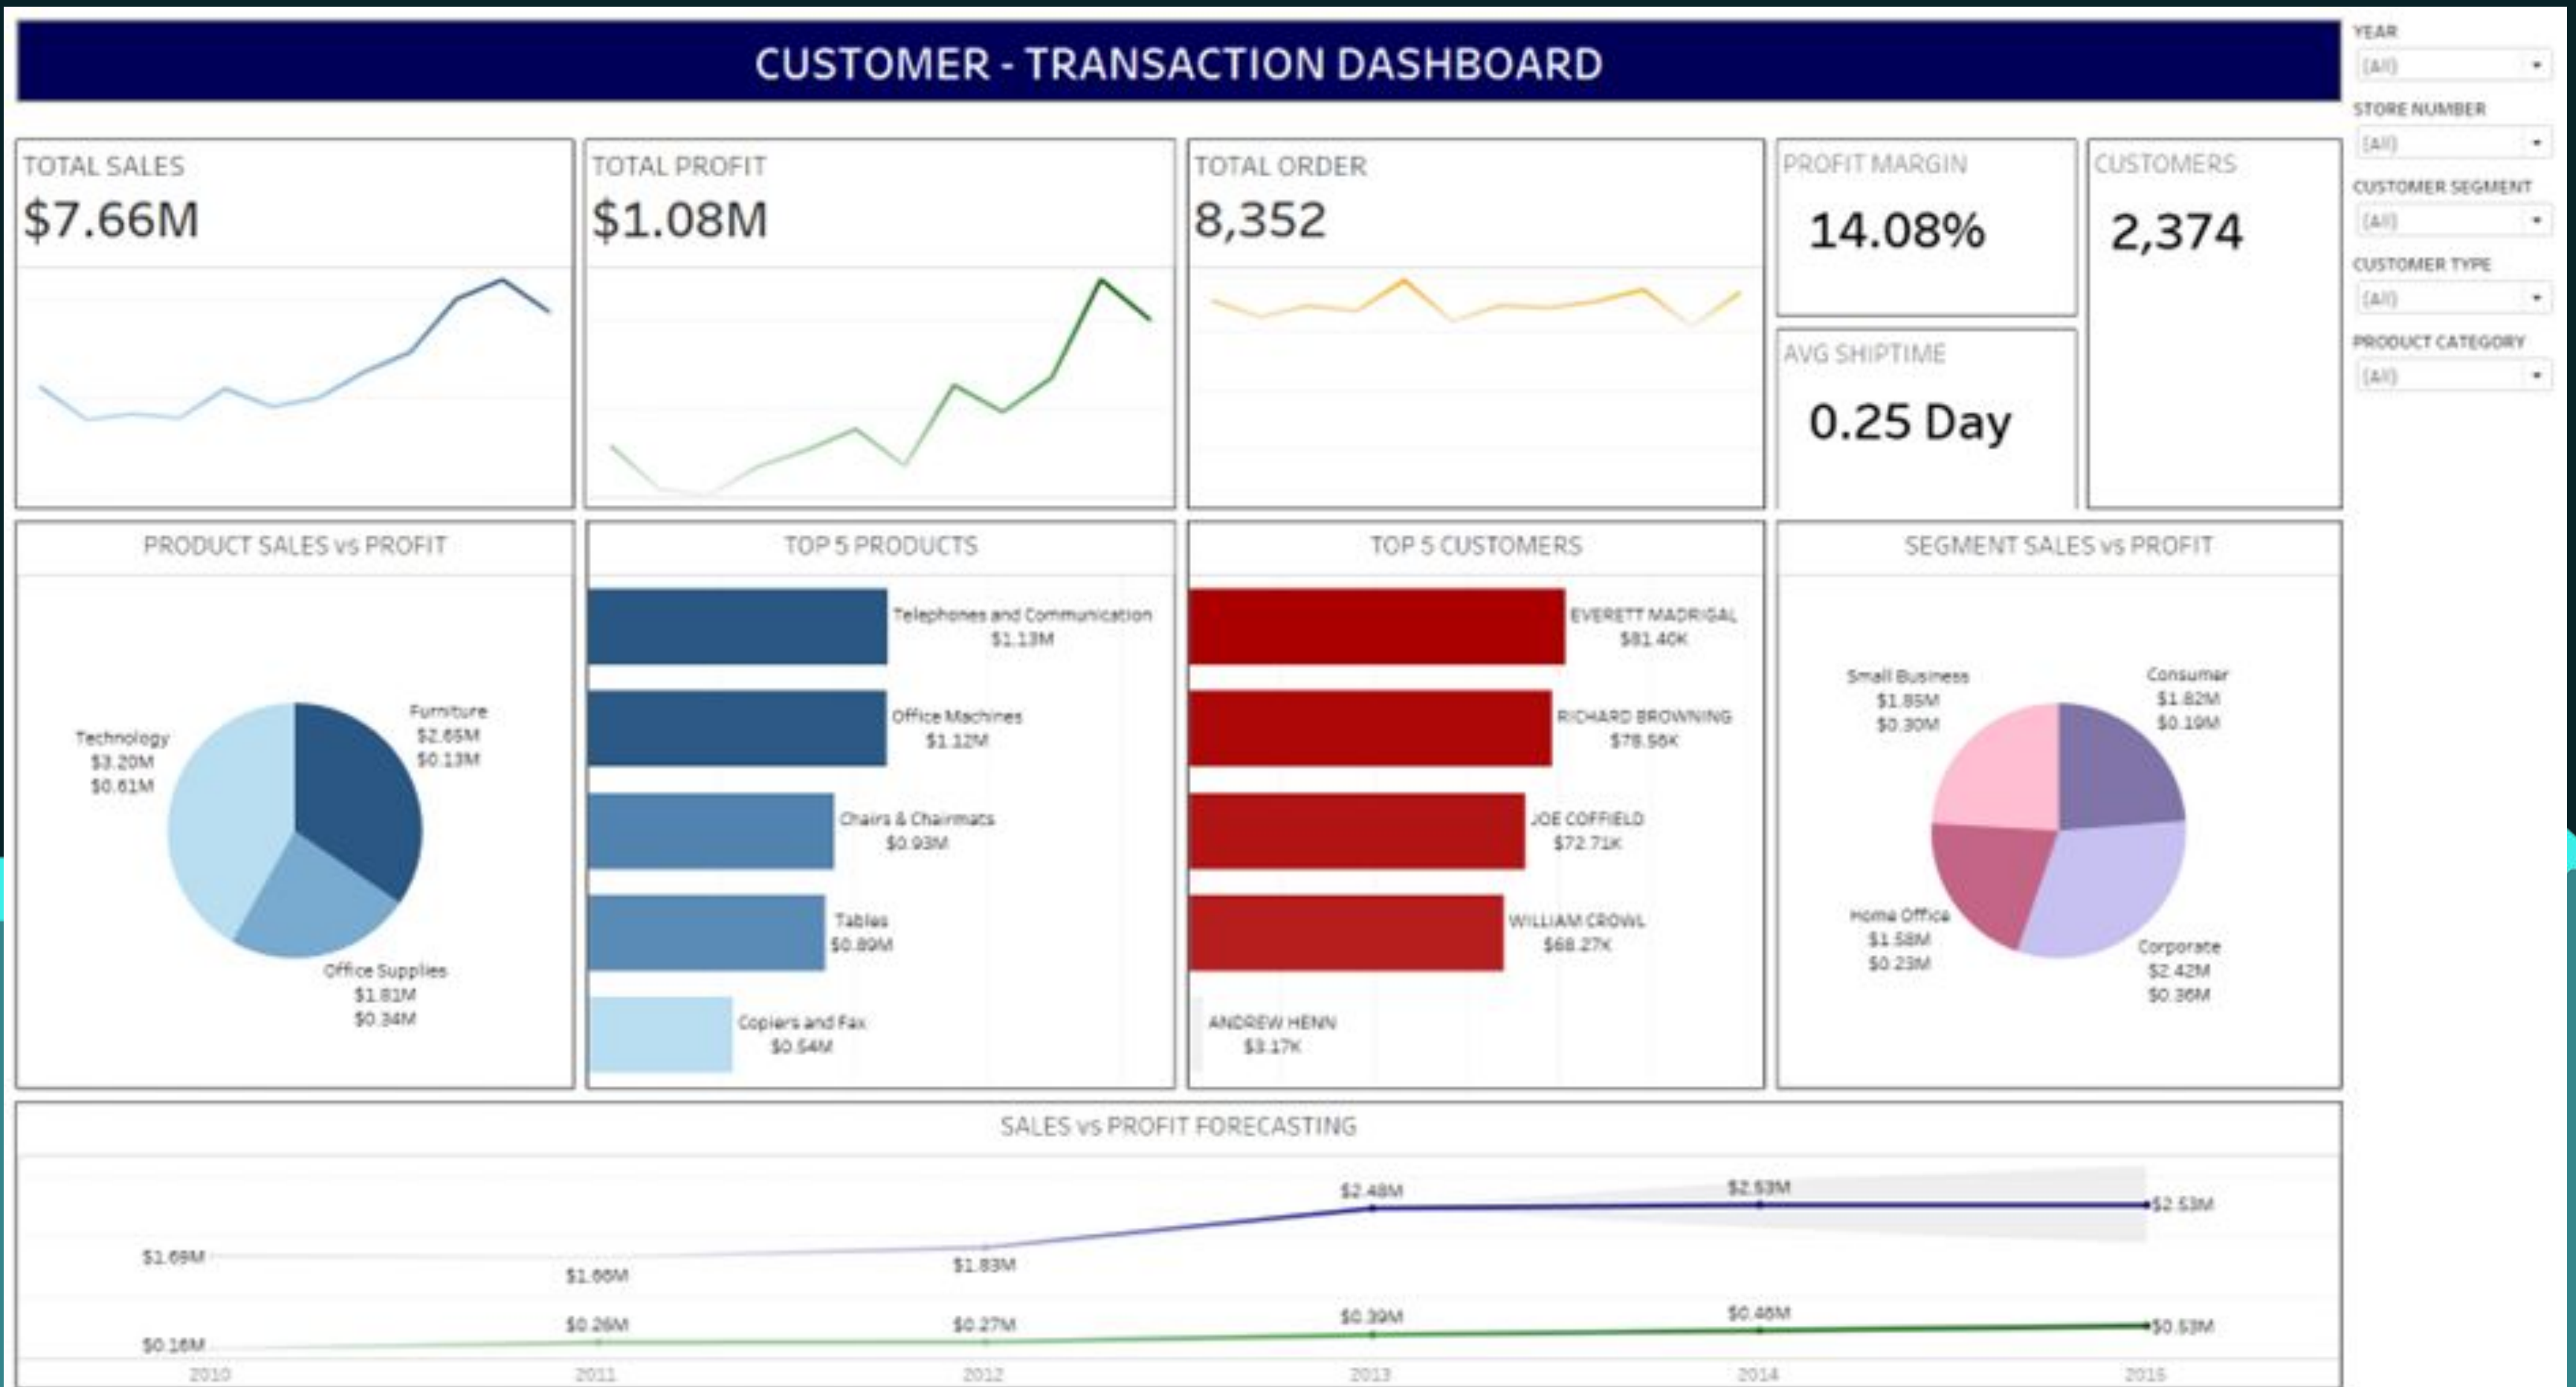
(All suitable features except Lat and Lon)

Classification Model	Accuracy
K-Nearest Neighbors	0.78372
Decision Tree	0.9573
Random Forest	0.91141
SVM (linear kernel)	0.78053
SVM (RBF kernel)	0.79489
SVM (sigmoid kernel)	0.69553
SVM (polynomial kernel)	0.7921
Neural Network Tensorflow	0.79409

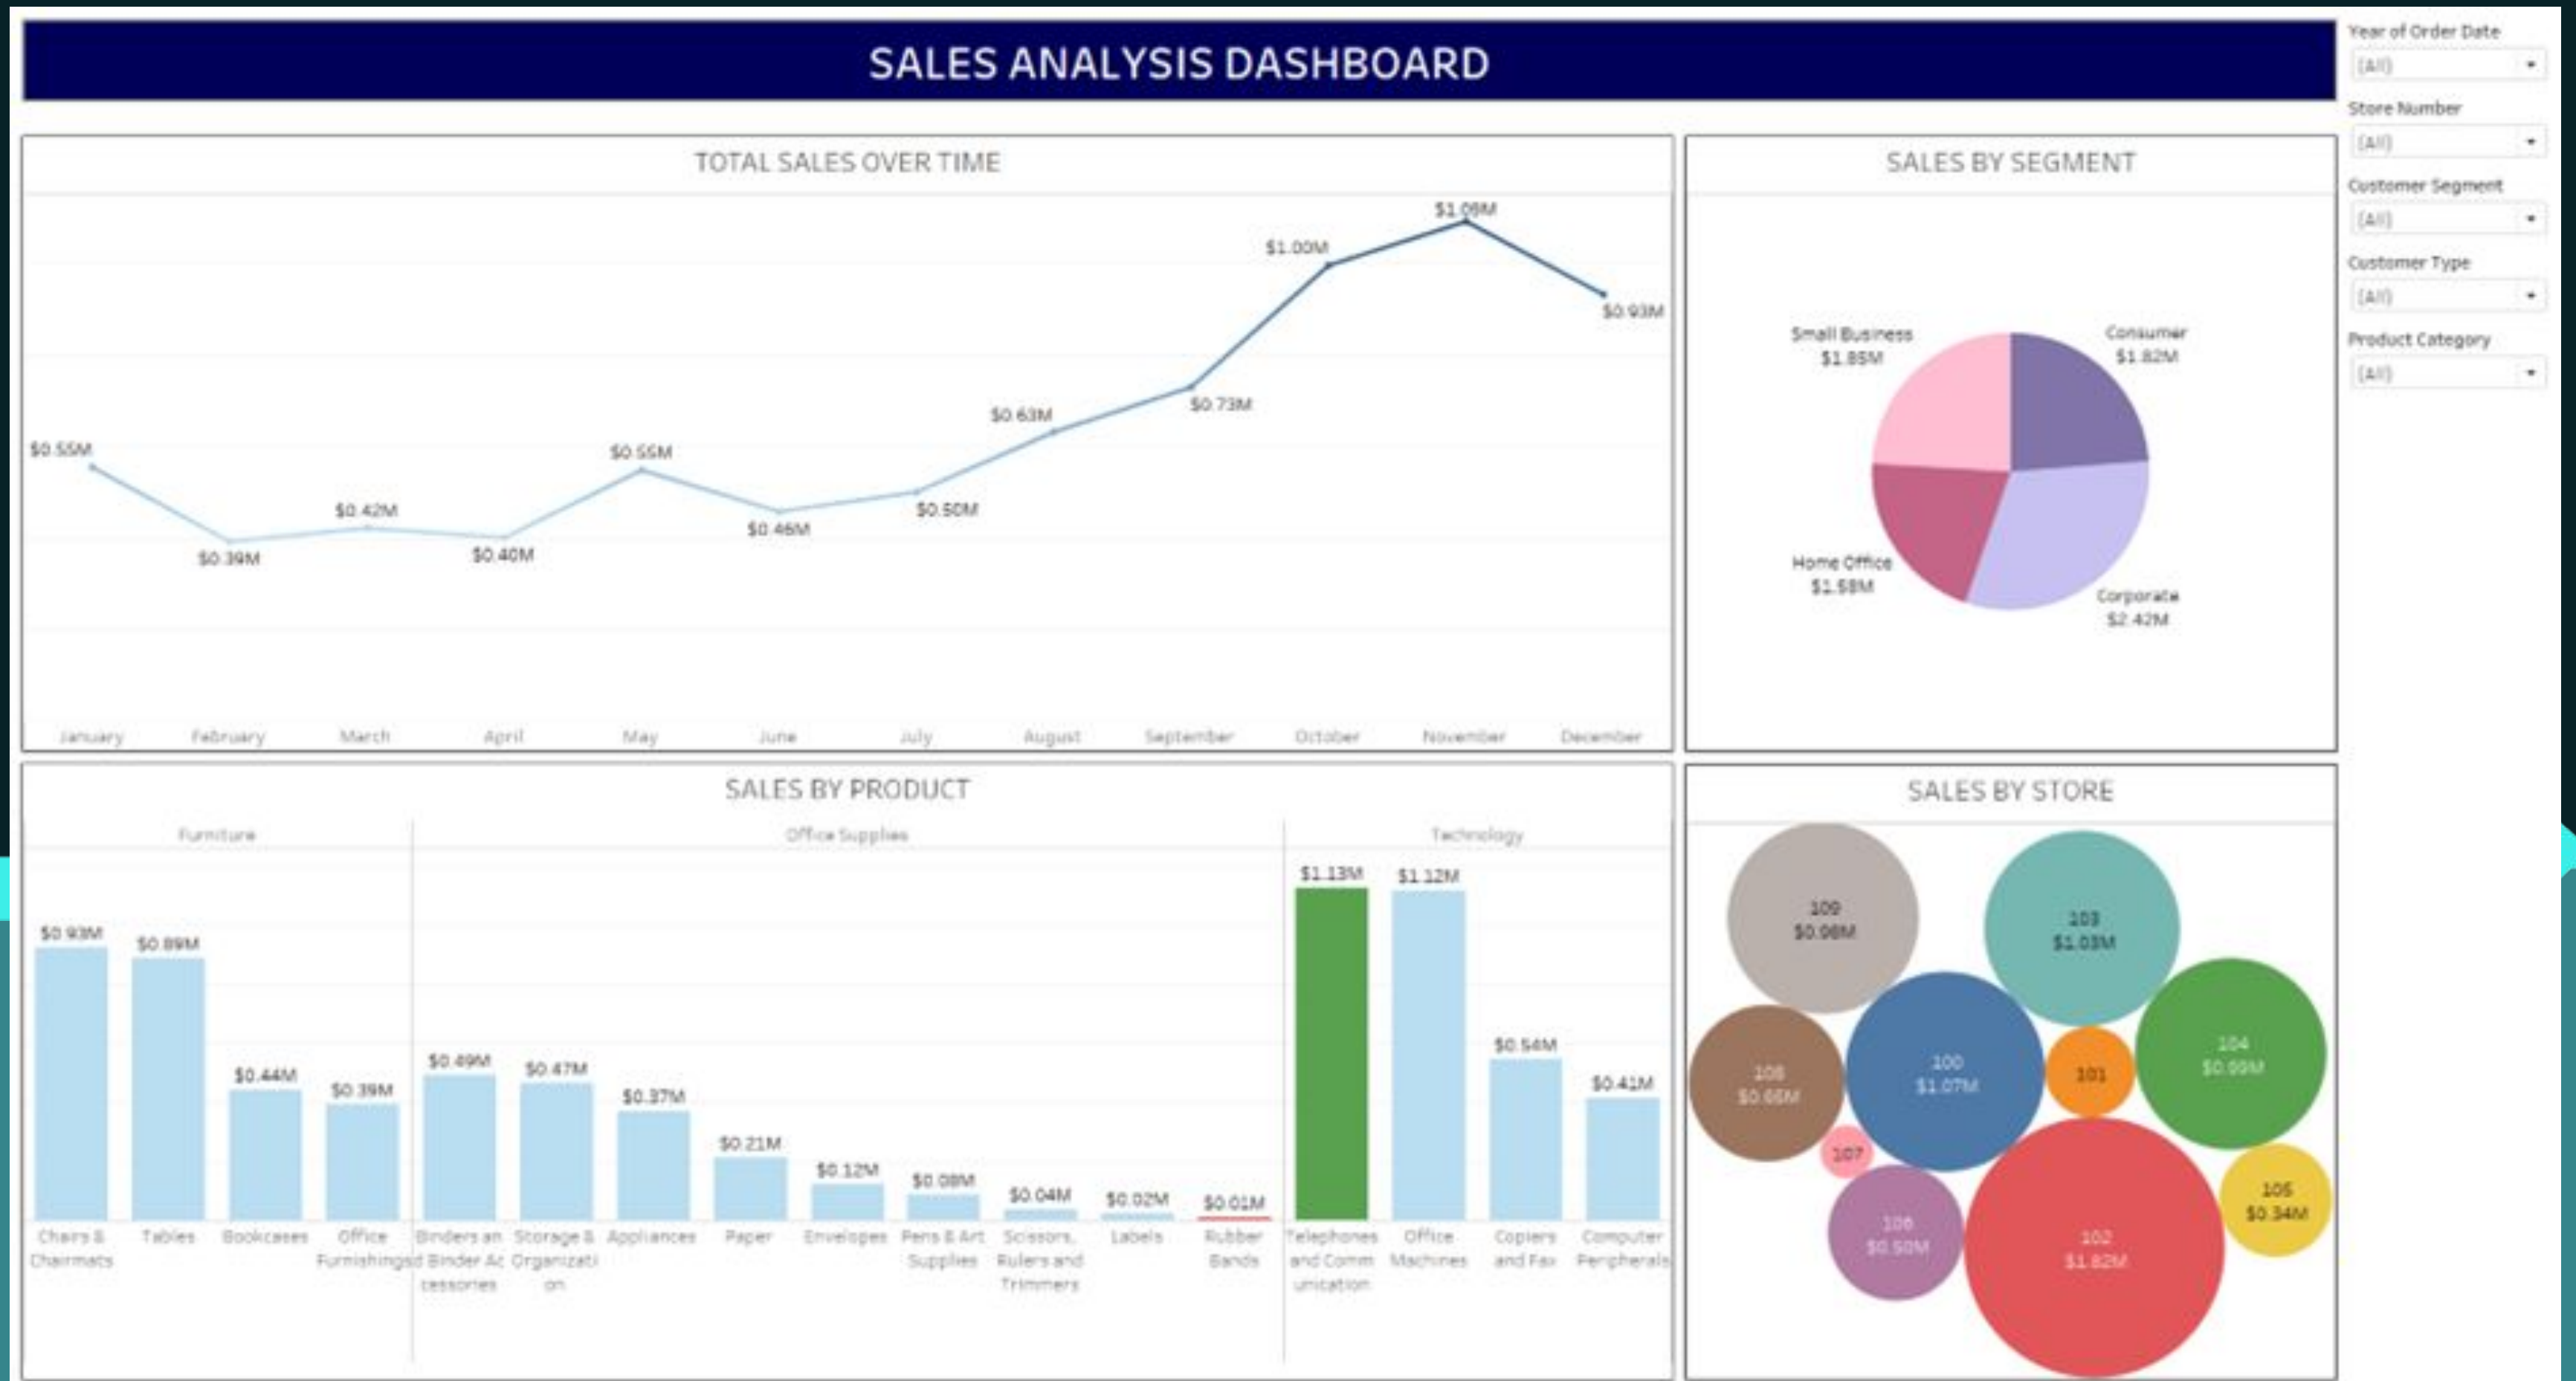
Best Model: Decision Tree (95%)

CONCLUSION: **RANDOM FOREST MODEL with FEATURE SET 1**

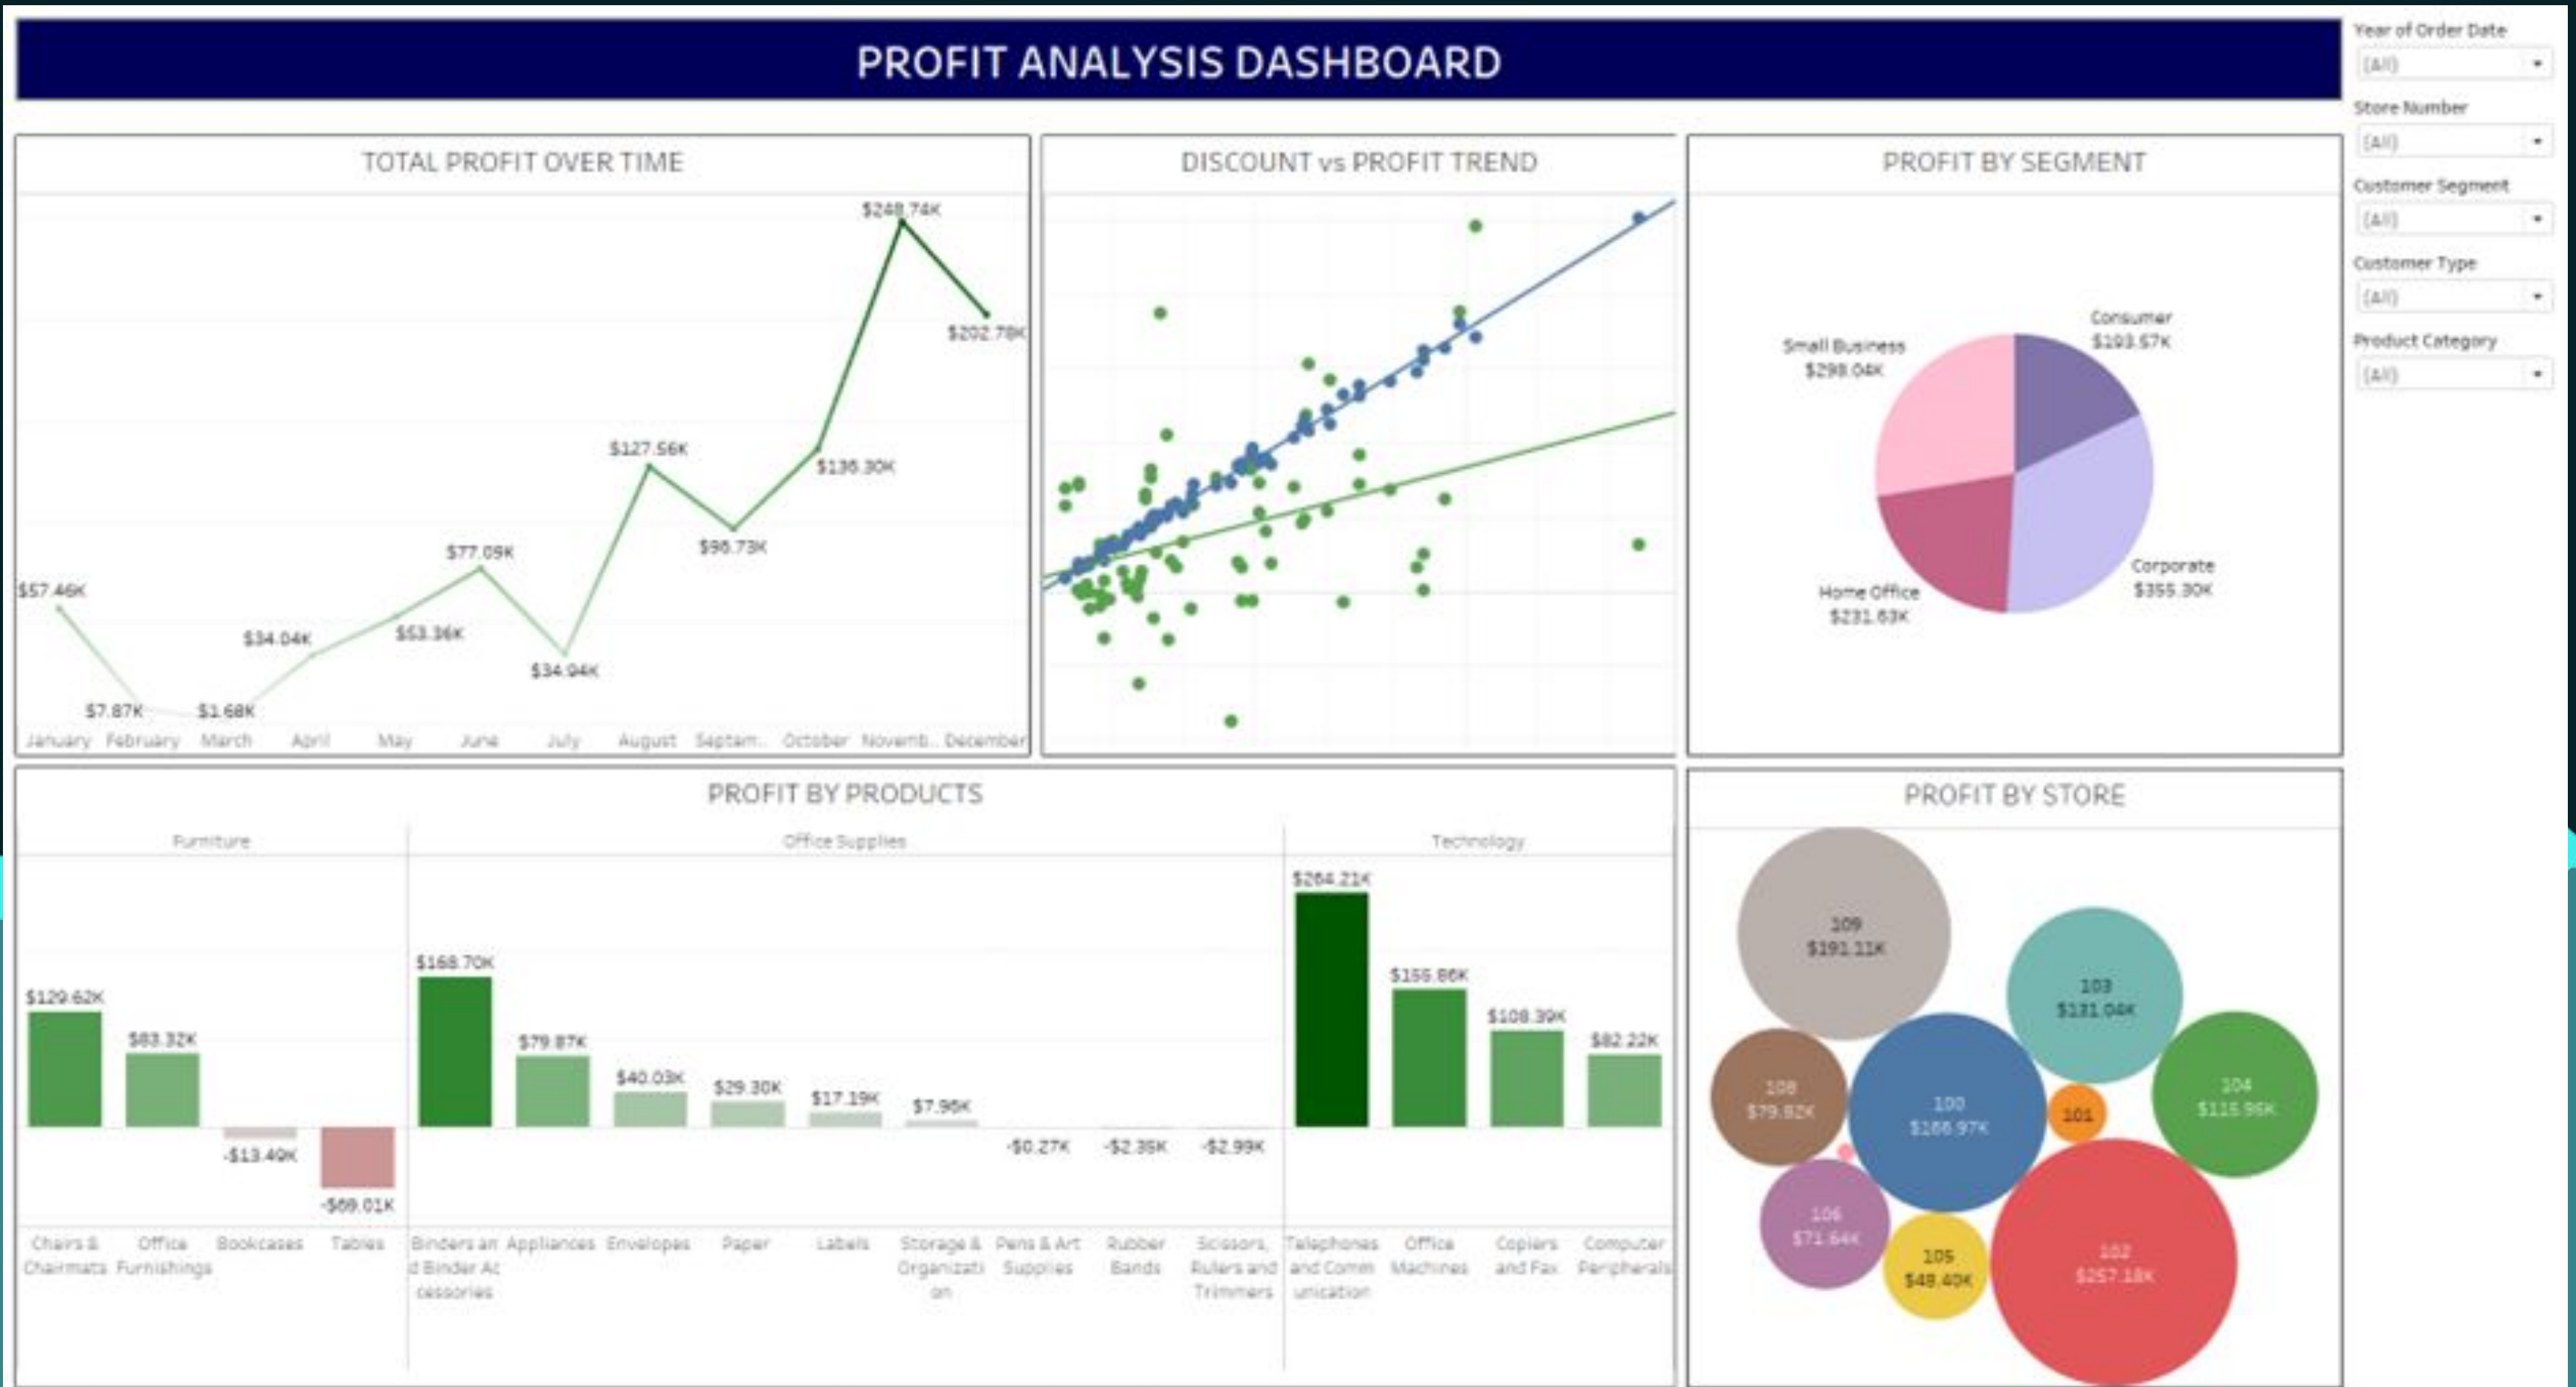
Visualization



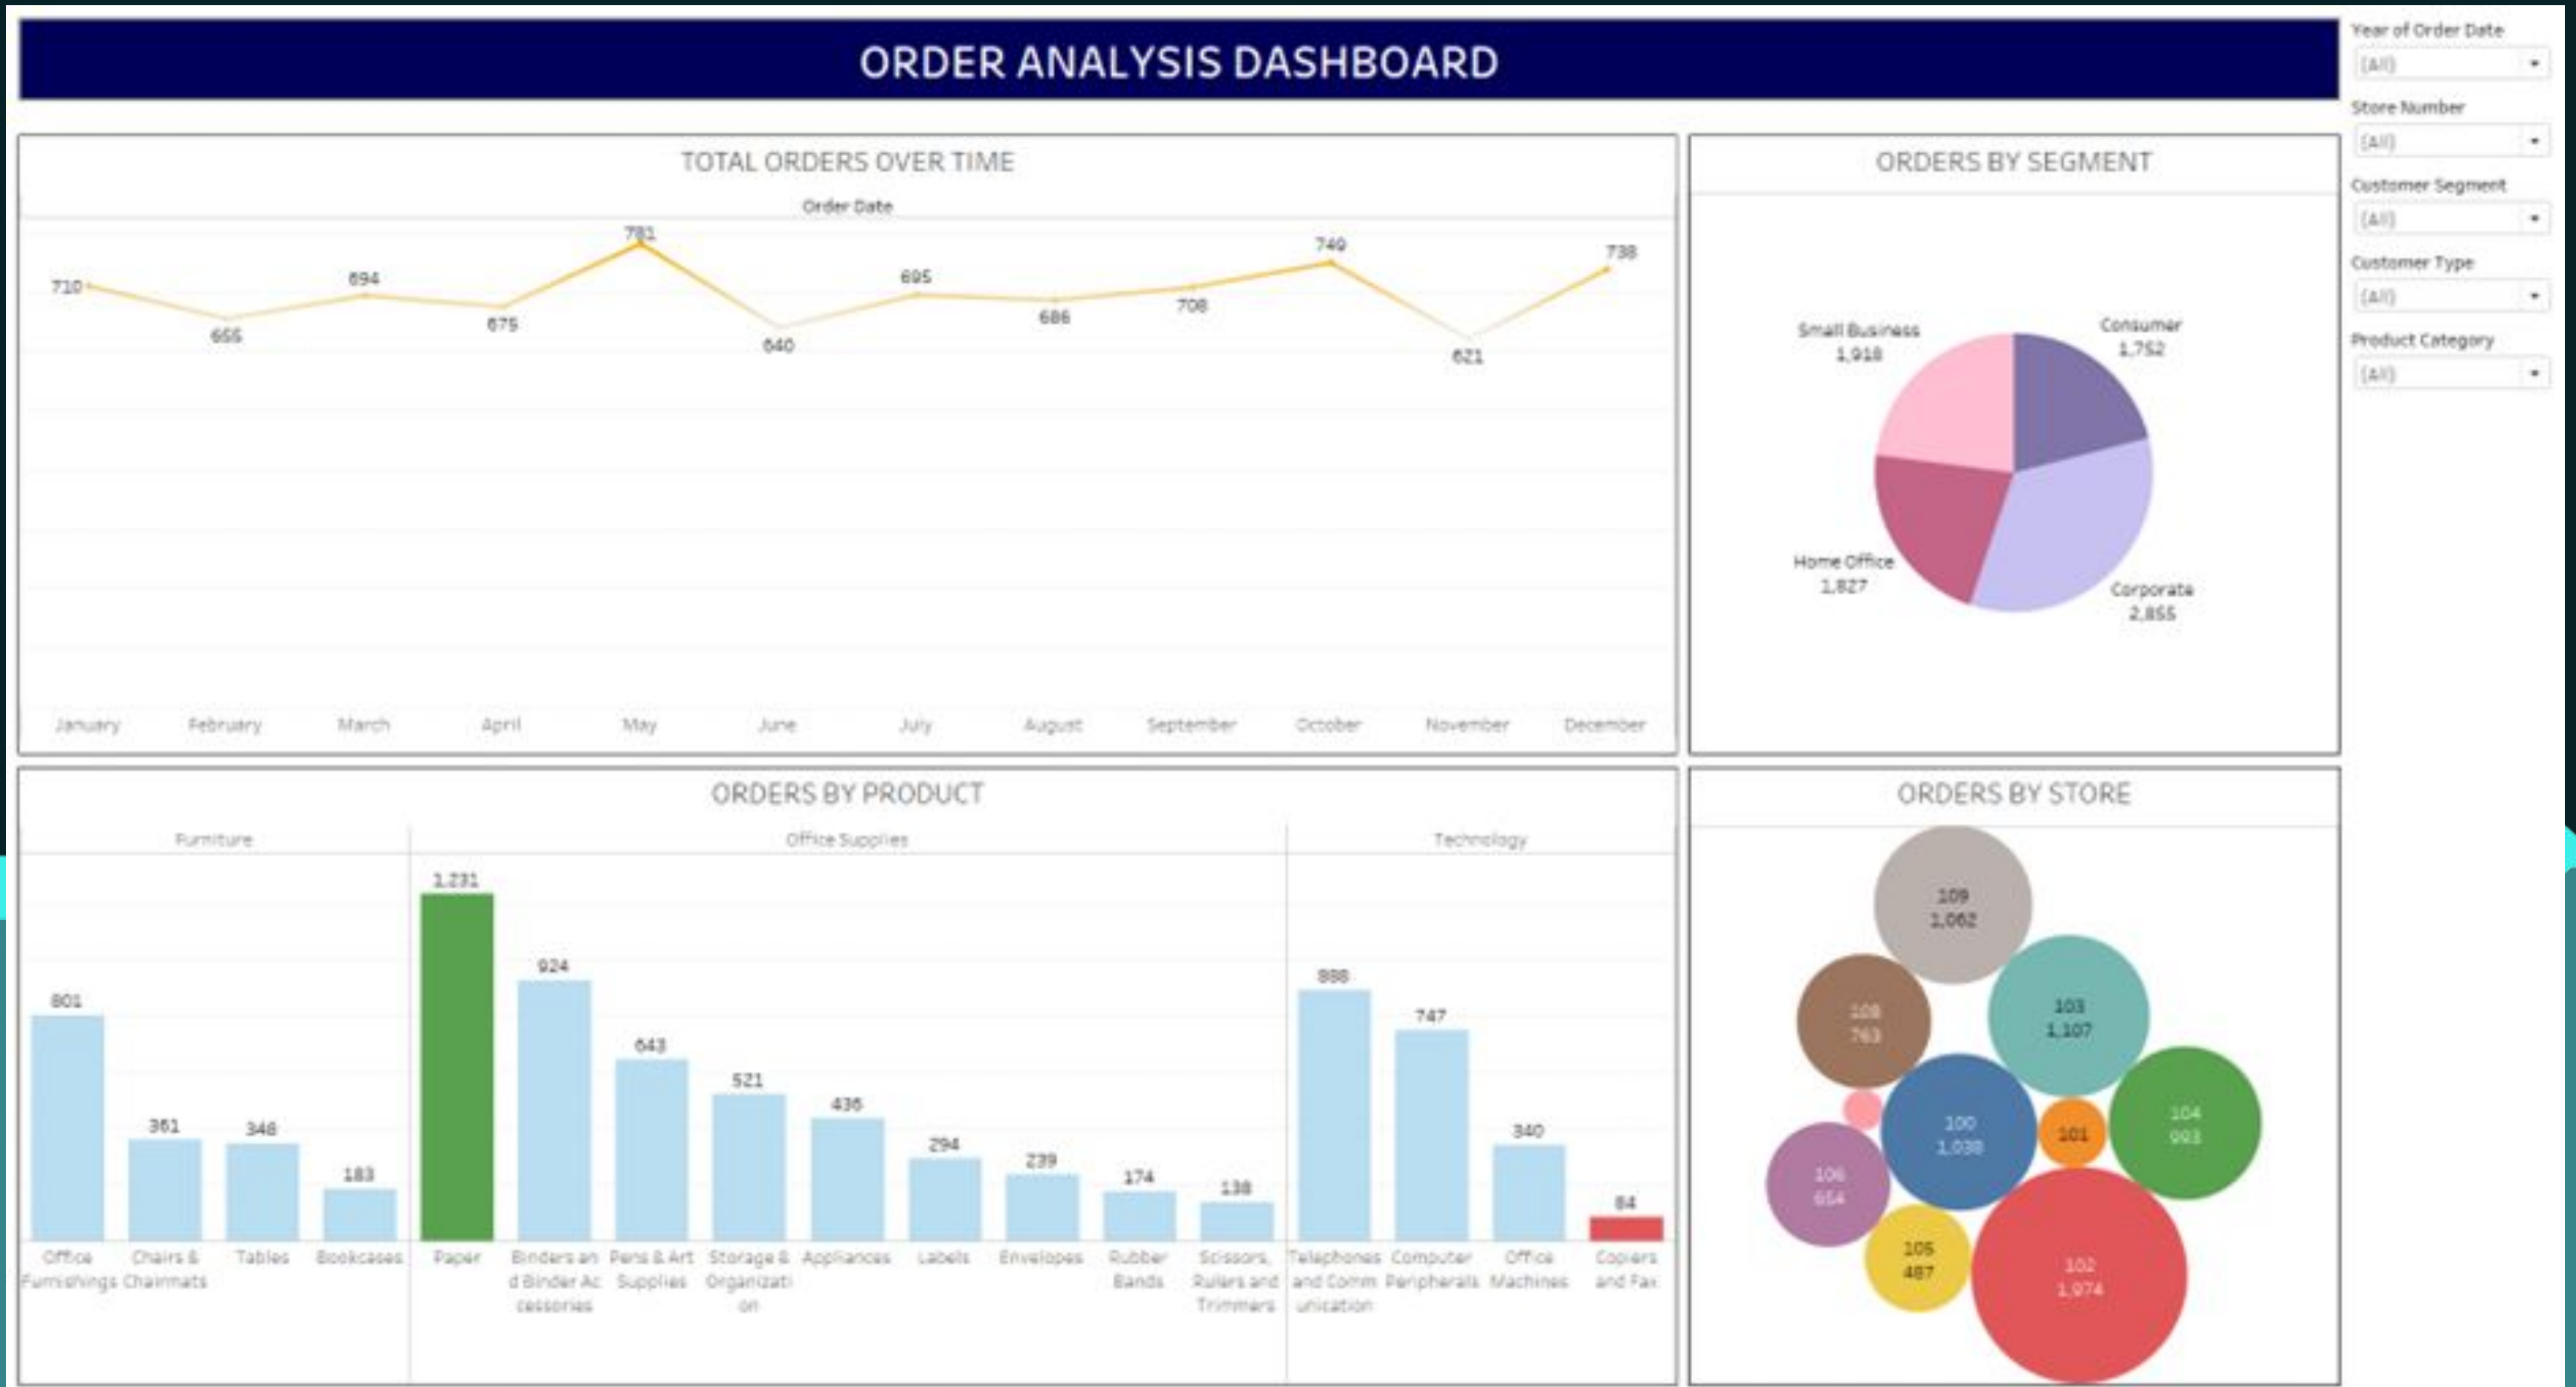
Visualization



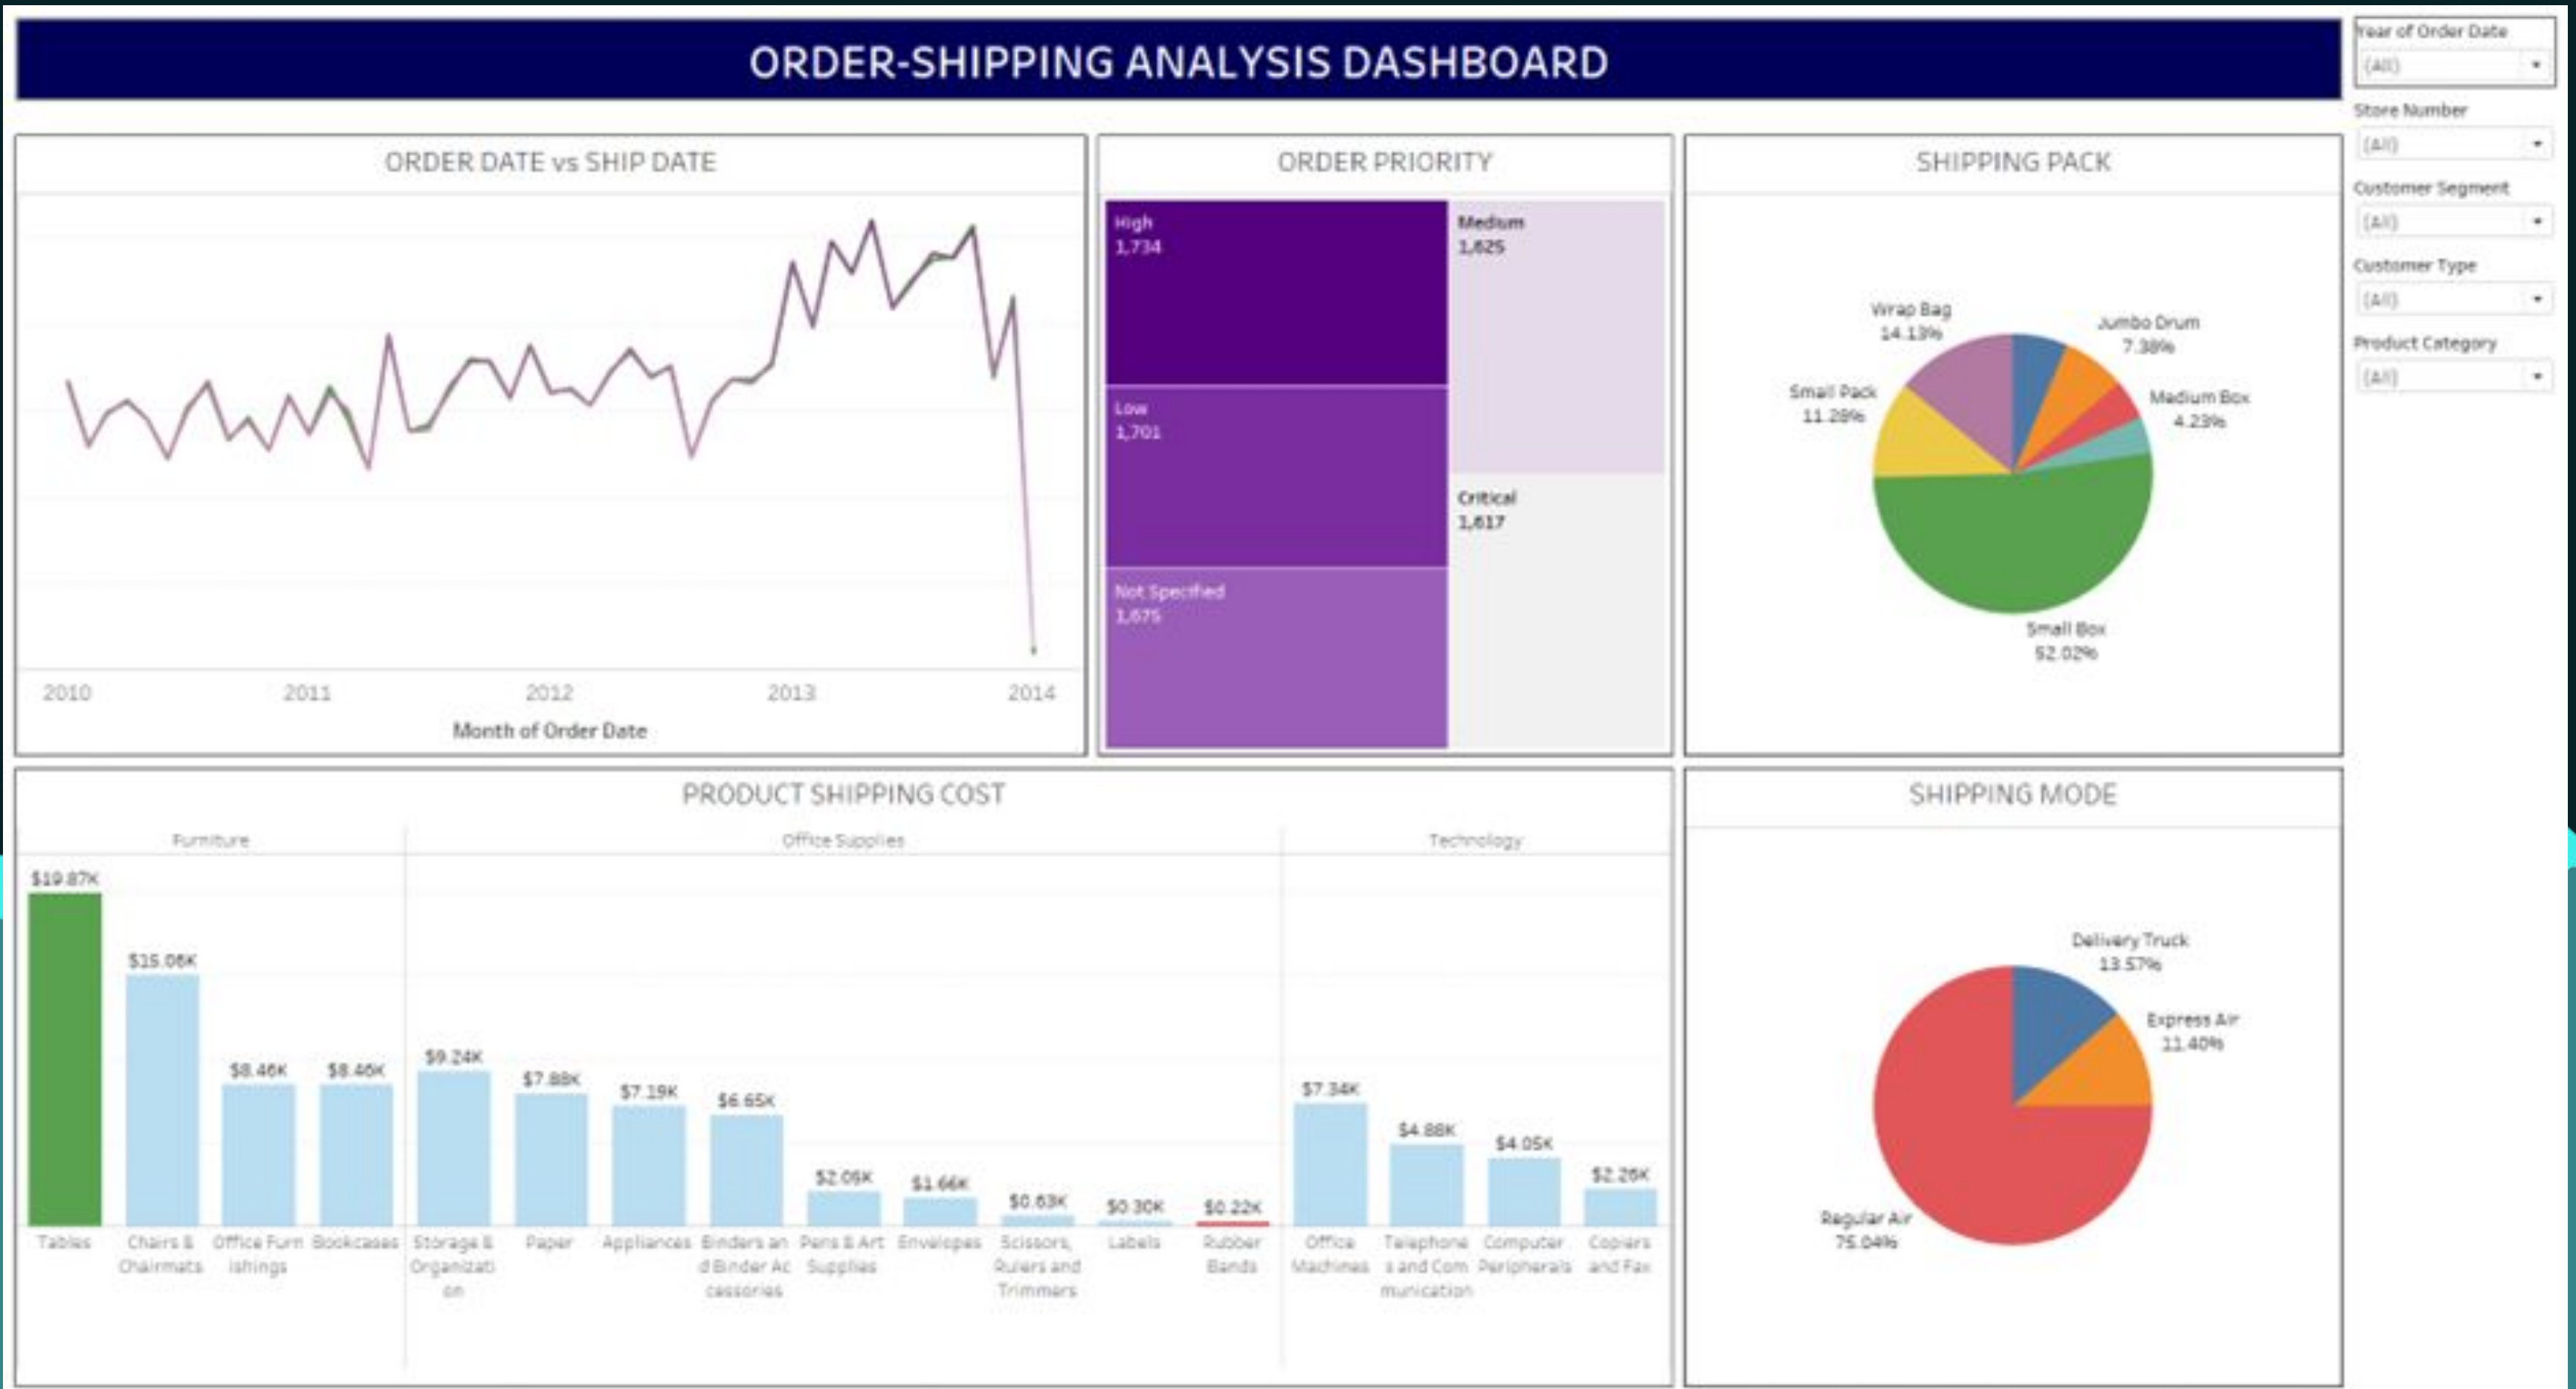
Visualization



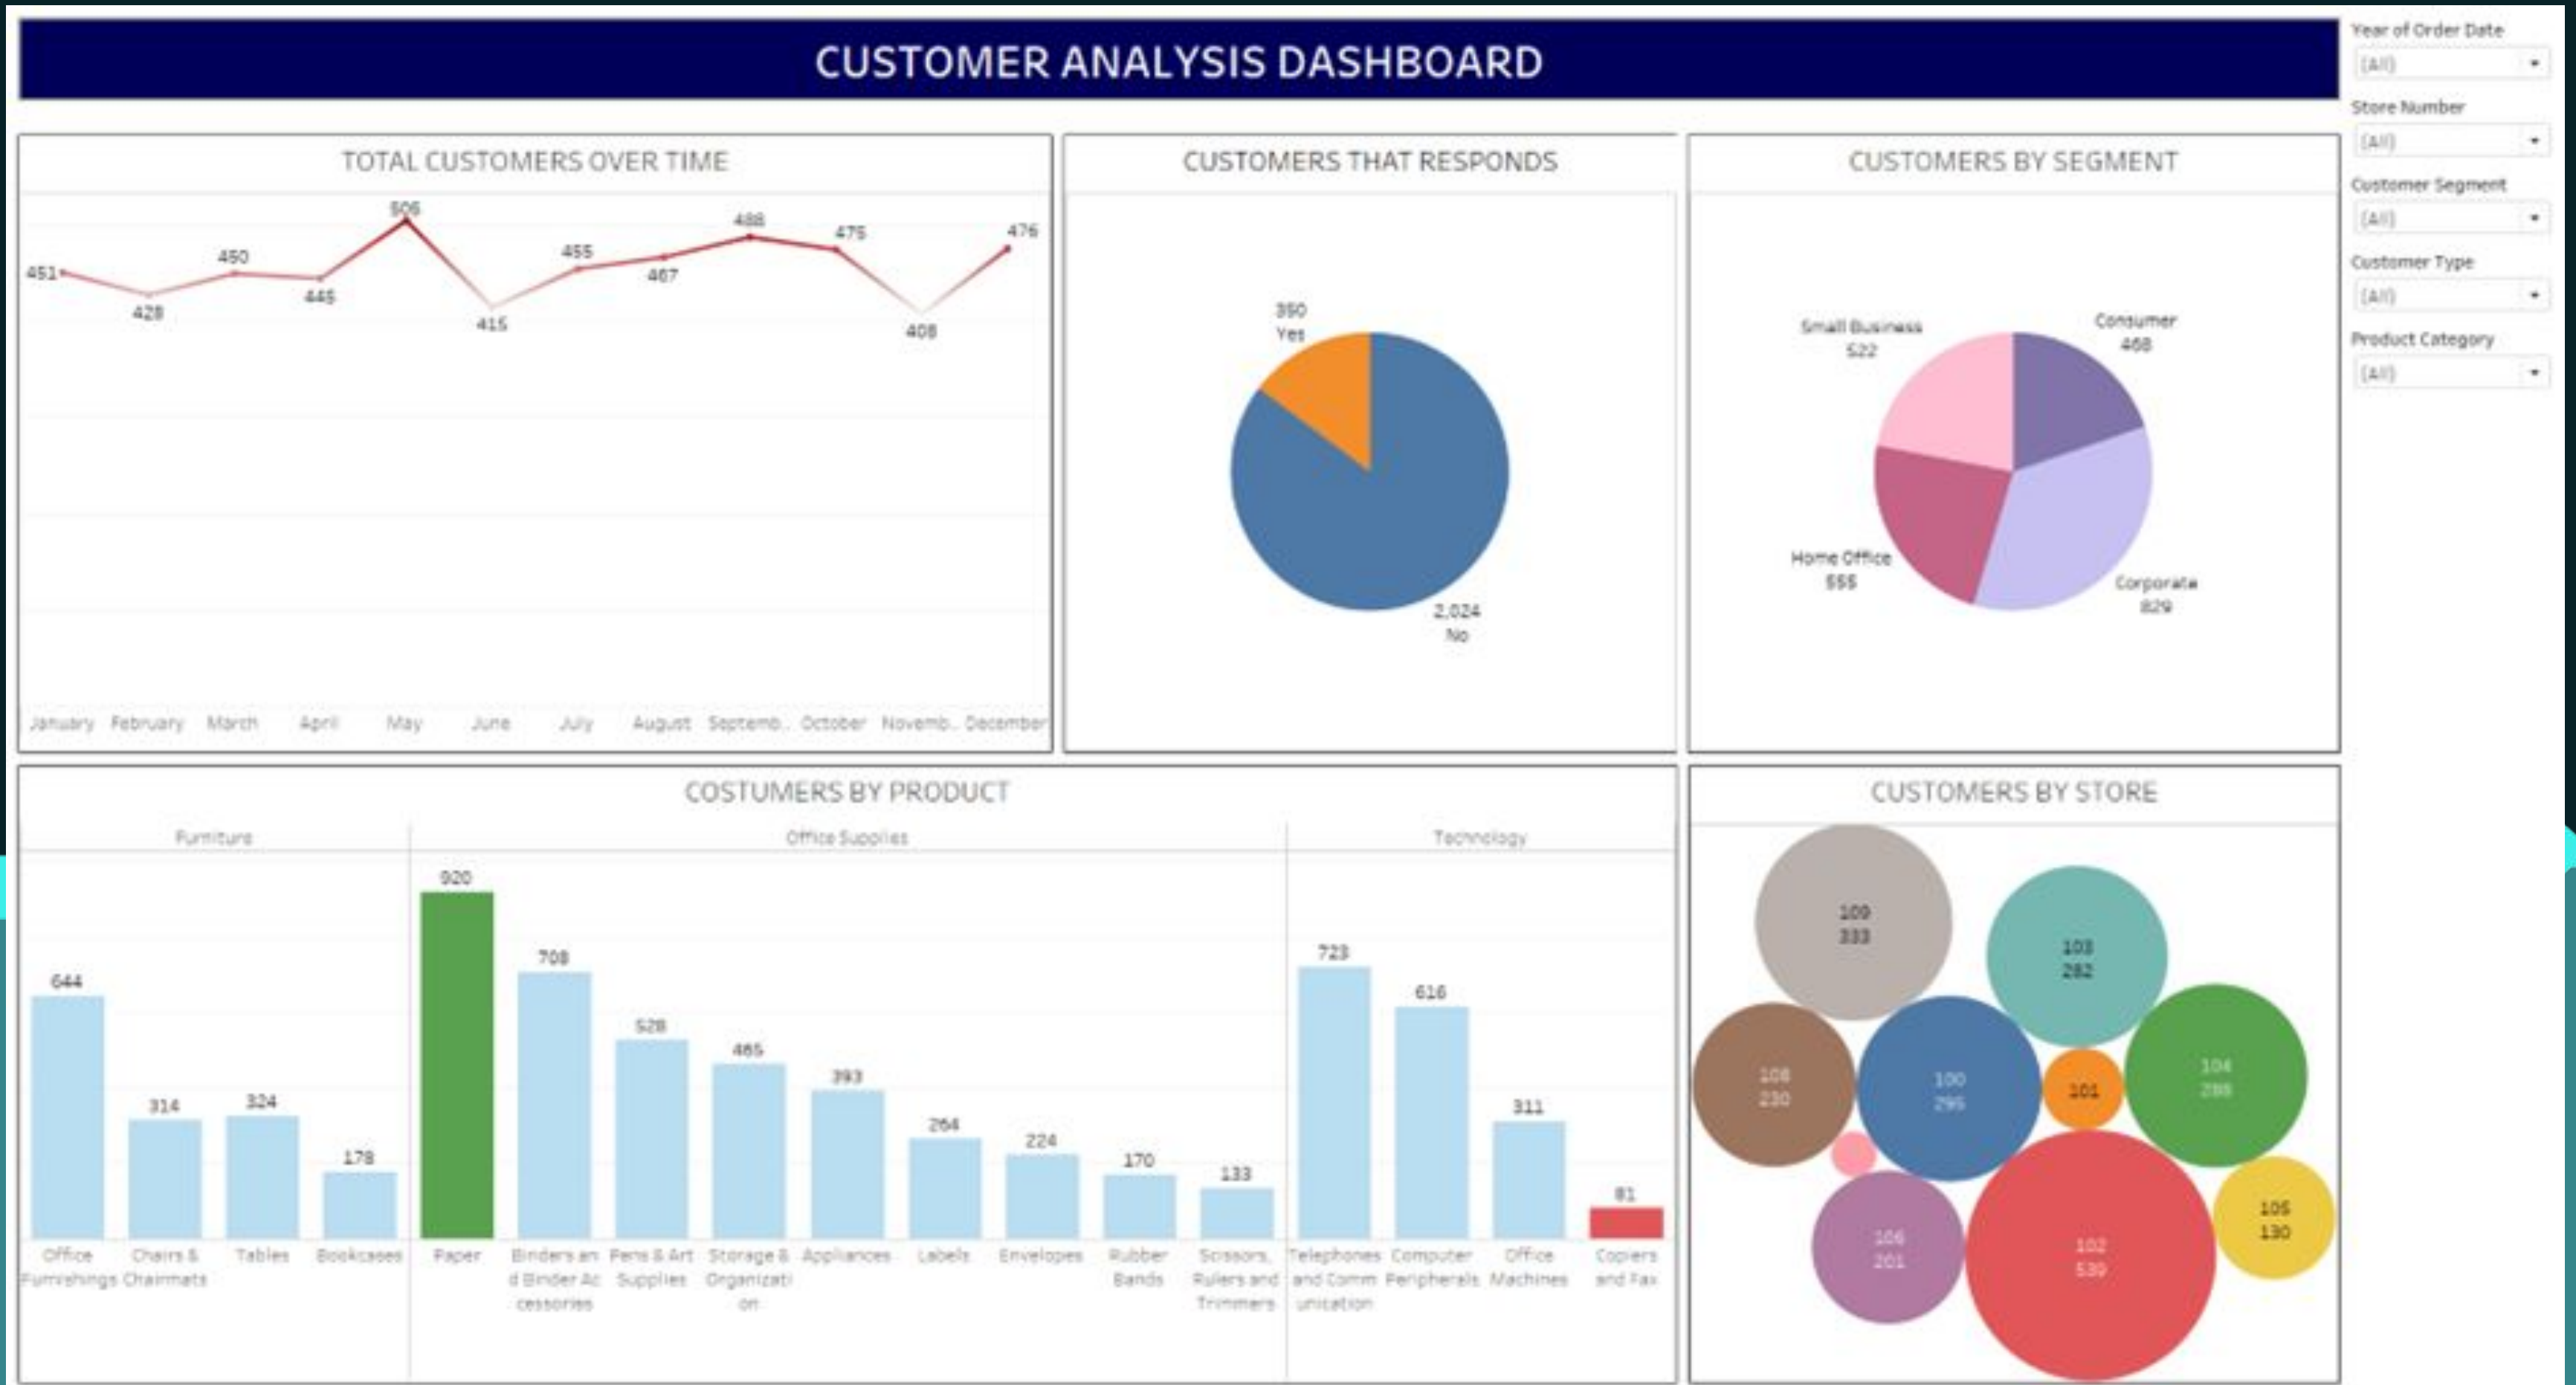
Visualization



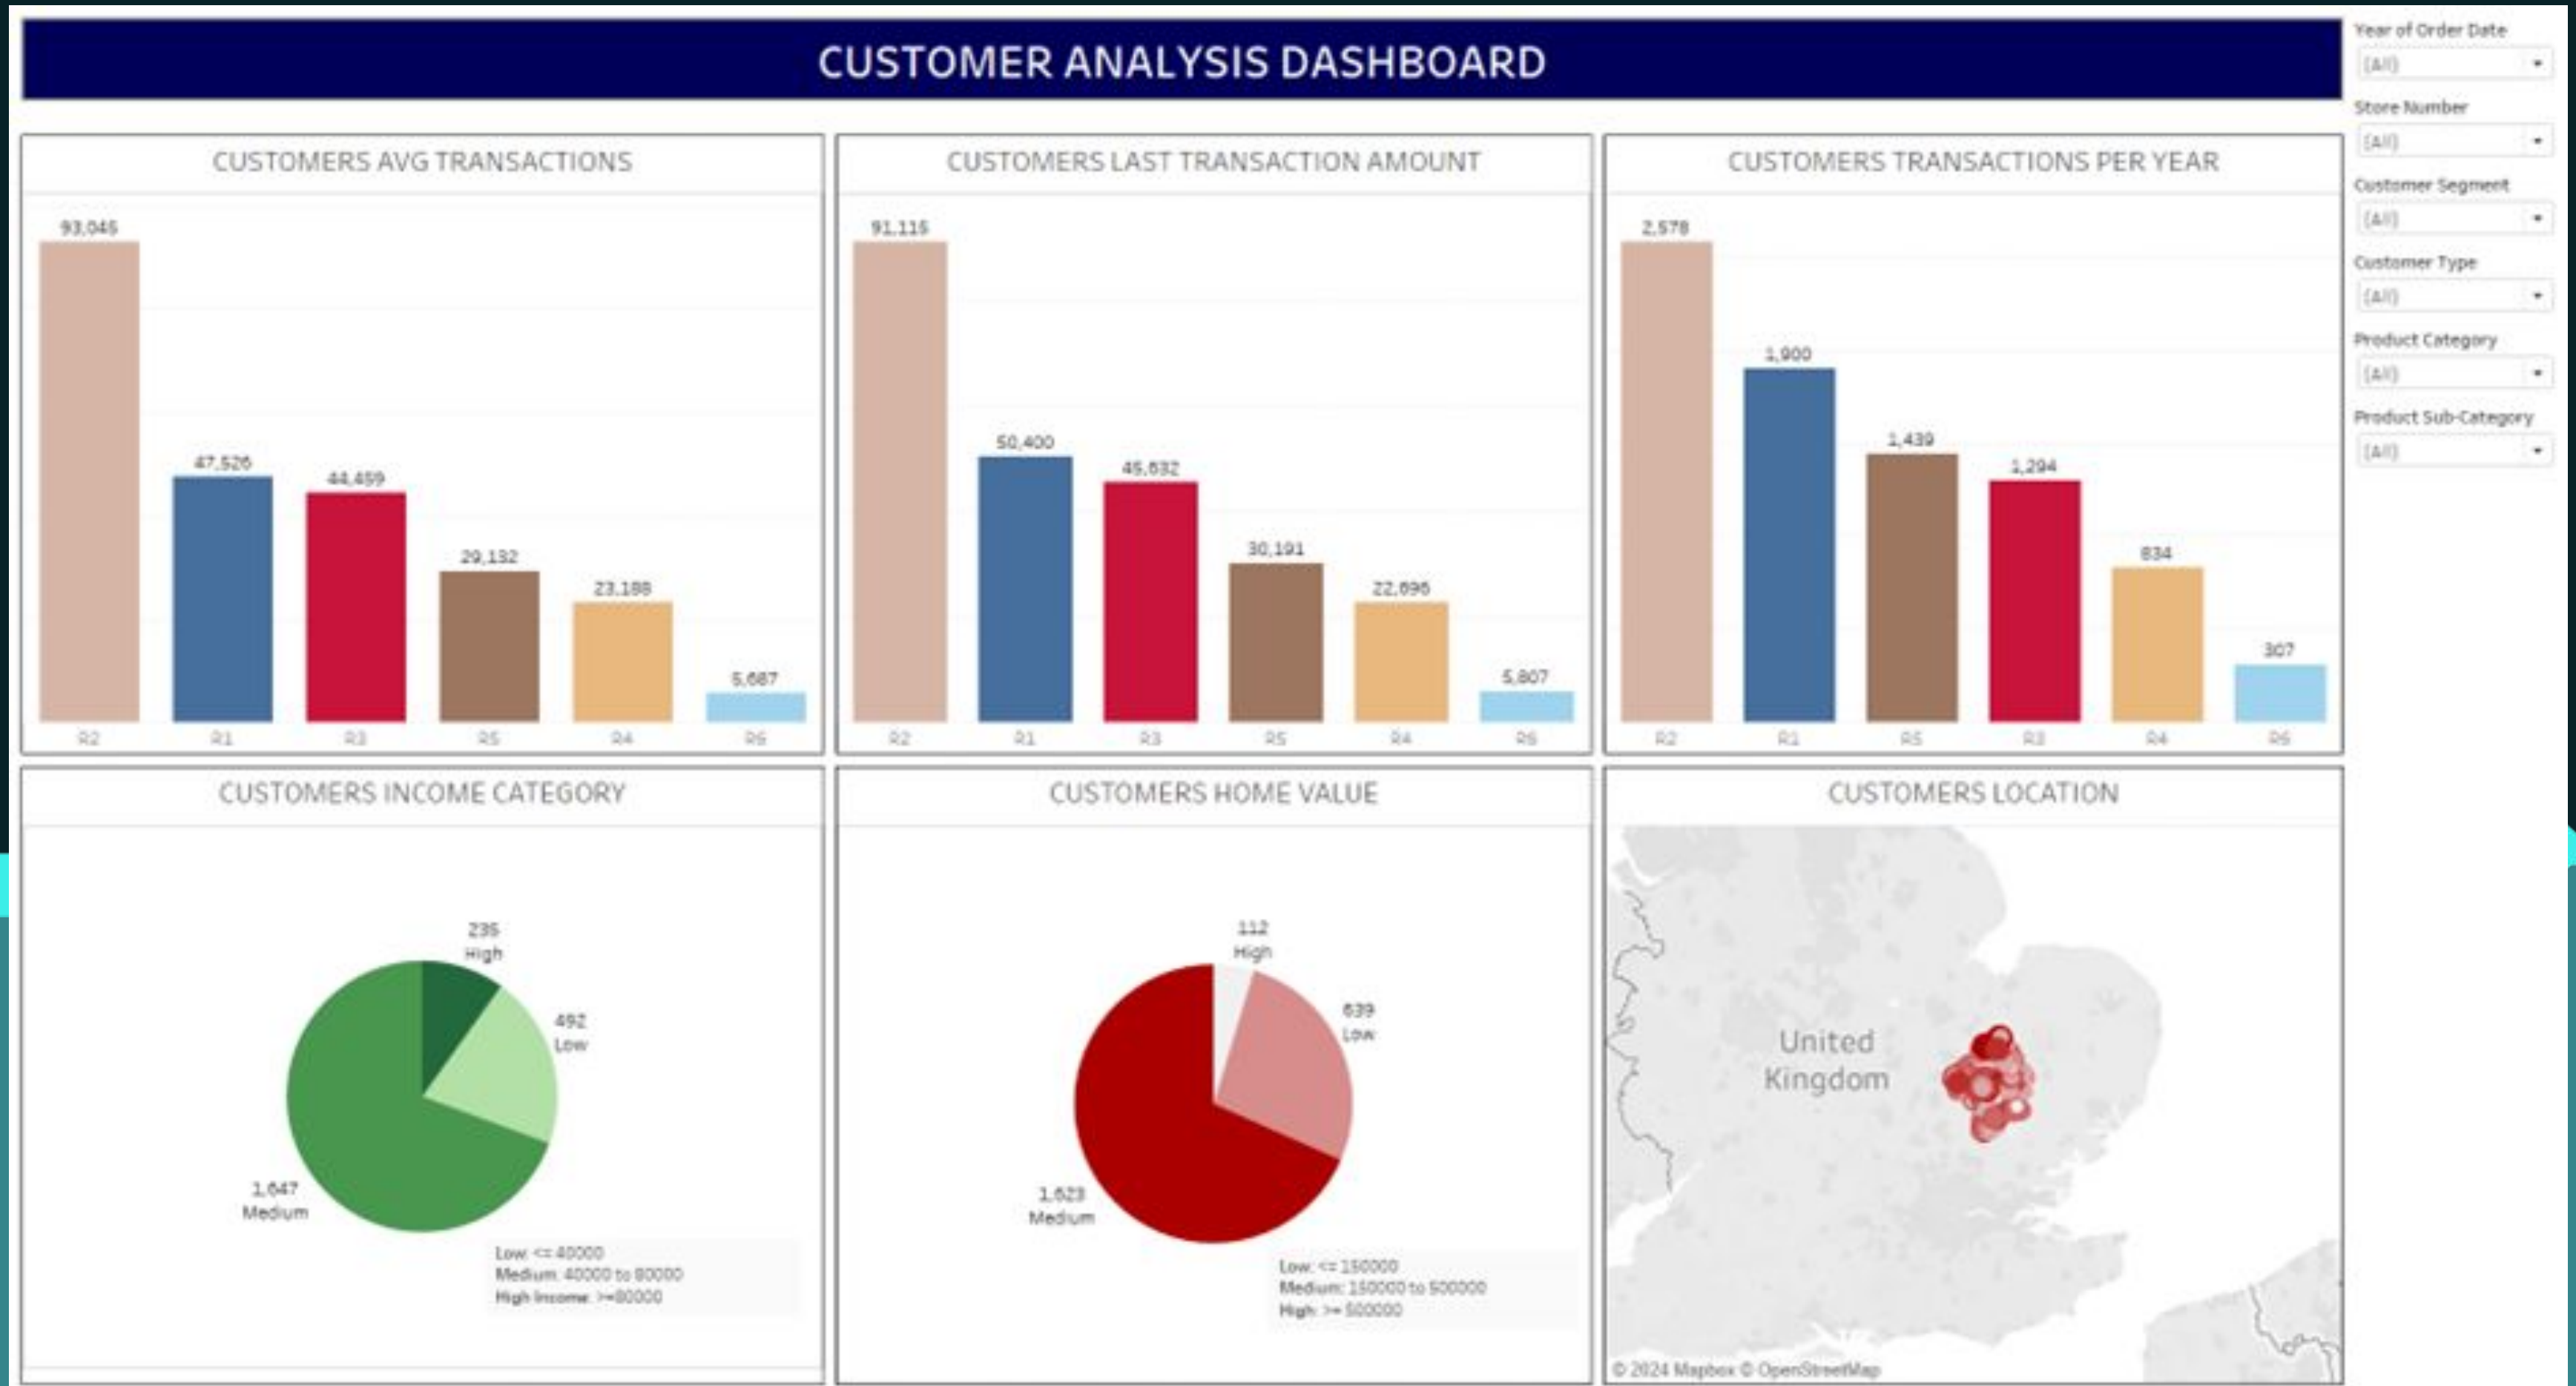
Visualization



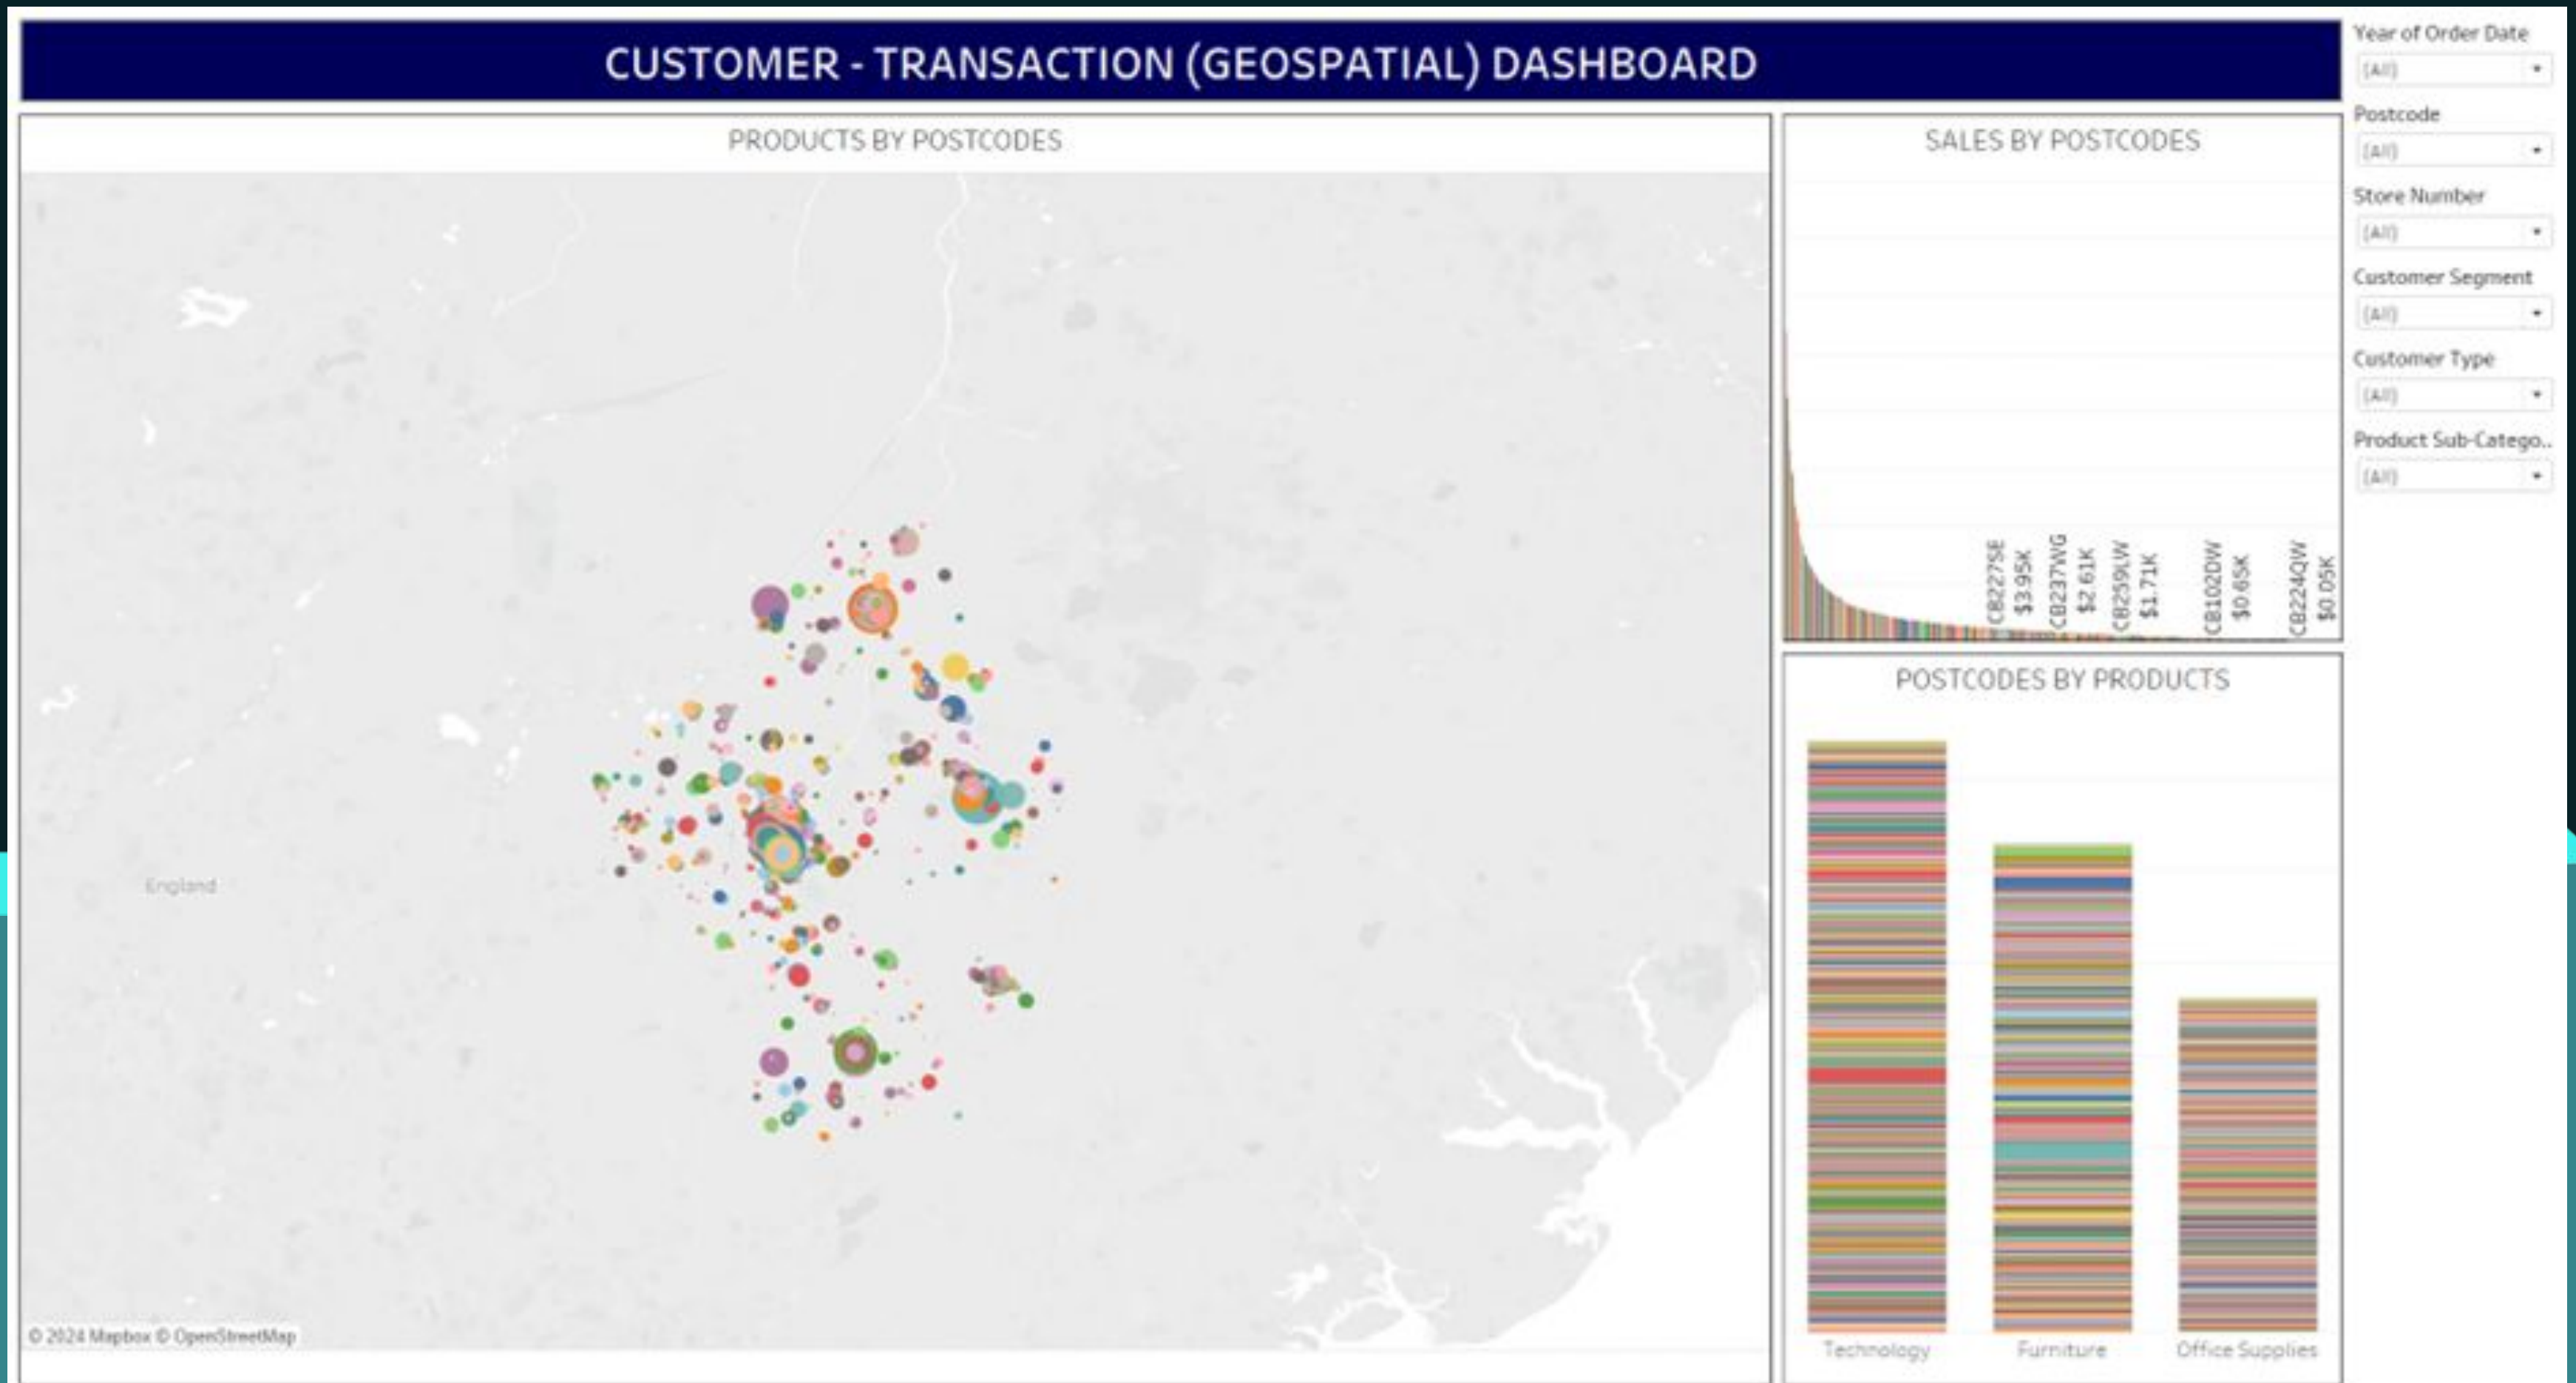
Visualization



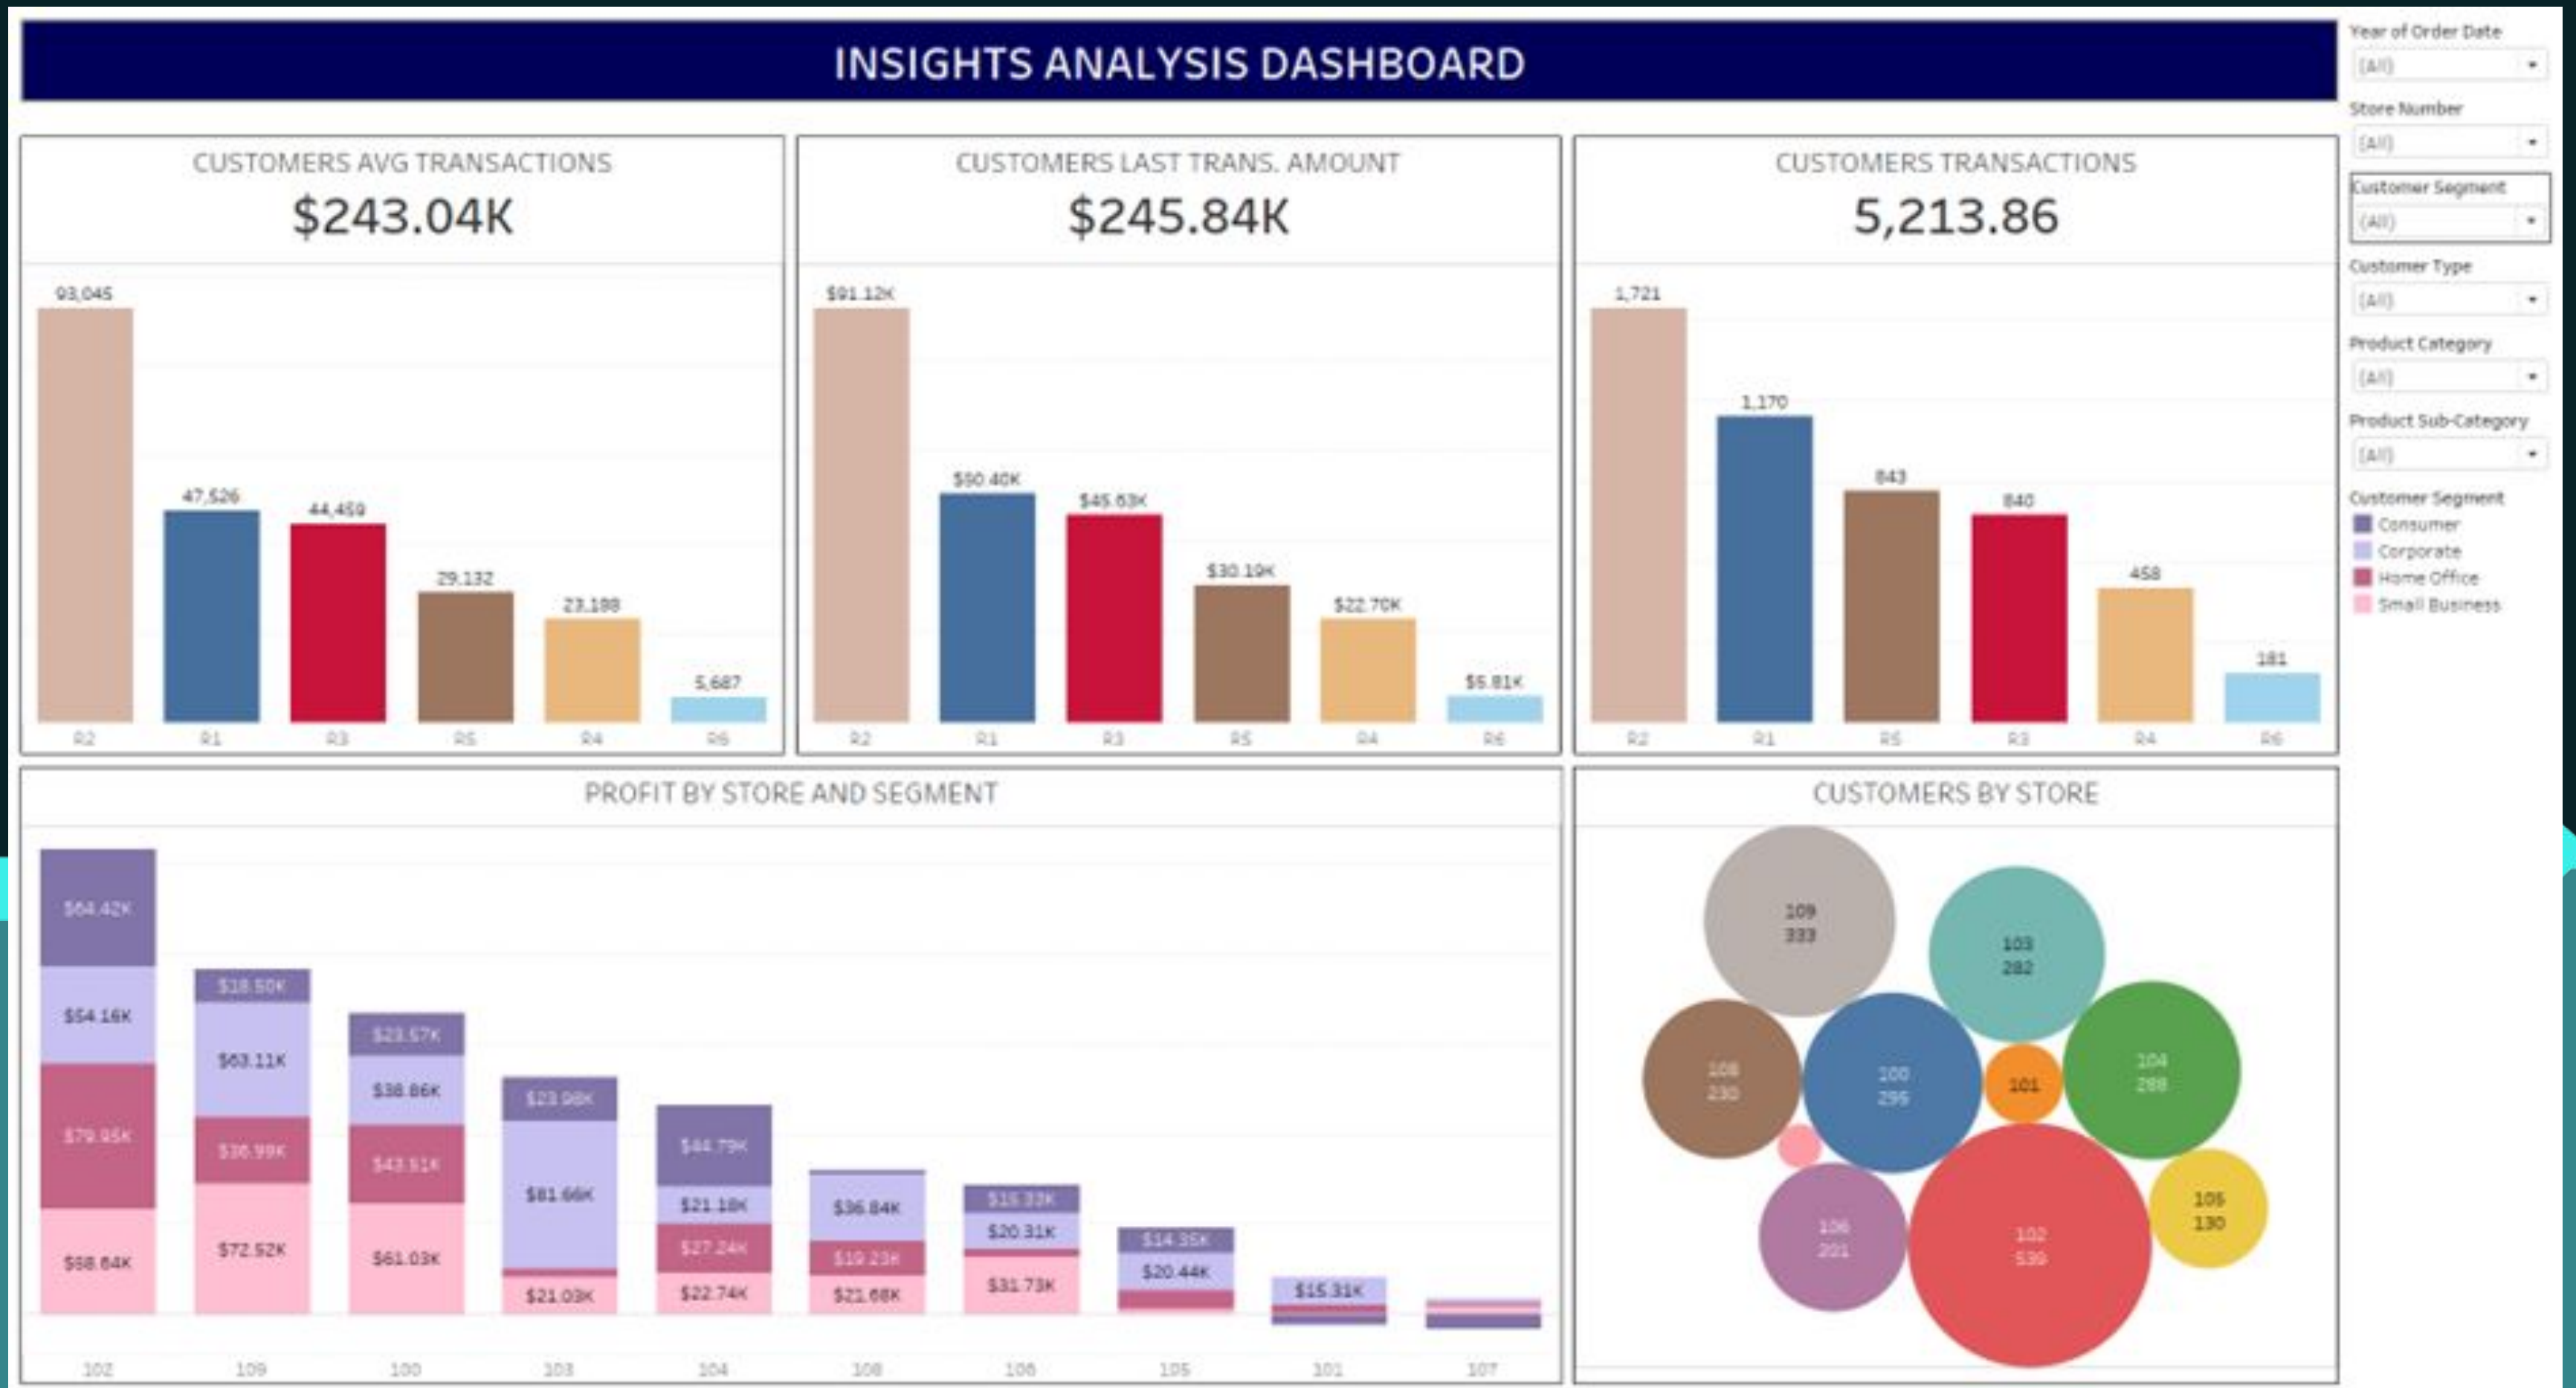
Visualization



Visualization



Visualization



Conclusion

This capstone project successfully leveraged the customer-transactions dataset by cleaning and preprocessing it, and using it to predict classification outcomes to provide valuable insights into patterns and trends that inform strategic decisions.

By developing dashboards in Tableau, the data was transformed into accessible, interactive visualizations that communicate vital findings effectively to stakeholders.

For the future, it is highly recommended to use a real dataset from a real industry and collect it in real-time instead of using dataset samples from the Internet, as well as explore additional visualization features in Tableau to engage users further.



THANK YOU

